

# Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Alexander Dean

Rochester Institute of Technology  
Golisano College of Computer and Information Sciences

August 12, 2014

# Process Cooperativity as a Feedback Metric

## in Concurrent Message-Passing Languages

### 1 Background

- Runtime Scheduling
- Cooperativity
- Message Passing

### 2 ErLam Toolkit

- The Language
- Channel Implementations
- Simulation & Visualization

### 3 Scheduler Implementations

- Example Schedulers
- Feedback Mechanisms

### 4 Results

### 5 Conclusions & Future Work

- ErLam Toolkit
- Cooperative Schedulers
- Cooperativity as a Metric

# Background

## 1 Background

- Runtime Scheduling
- Cooperativity
- Message Passing

# Background: Runtime Scheduling

- Schedulers can be defined in a discrete manner:
  - 1 *Choose* a process from set,
  - 2 *Reduce* it,
  - 3 *Update* private scheduler state.

# Background: Runtime Scheduling

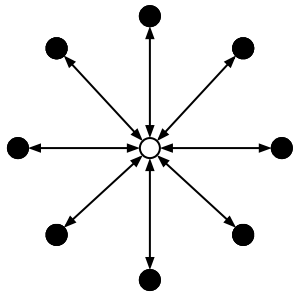
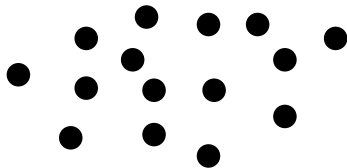
- Schedulers can be defined in a discrete manner:
  - 1 *Choose* a process from set,
  - 2 *Reduce* it,
  - 3 *Update* private scheduler state.
- Statistics can be gathered at every step about process:
  - Timestamp of last run,
  - Number of reductions, *etc.*

# Background: Runtime Scheduling

- Schedulers can be defined in a discrete manner:
  - 1 *Choose* a process from set,
  - 2 *Reduce* it,
  - 3 *Update* private scheduler state.
- Statistics can be gathered at every step about process:
  - Timestamp of last run,
  - Number of reductions, *etc.*
- *What statistics are useful?*

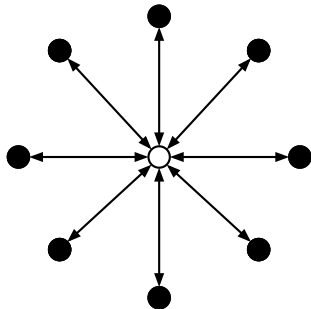
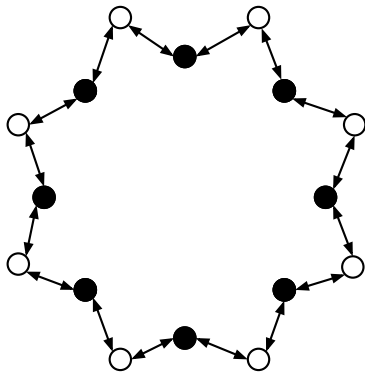
# Background: Cooperativity

What is Process Cooperativity?



# Background: Cooperativity

What does Cooperativity give us?





# Background: Message Passing

We use a Symmetric, Synchronous, Message-Passing Primitive:

swap

- Purely captures cooperation of processes through synchronizing on a shared channel.

## 2 ErLam Toolkit

- The Language
- Channel Implementations
- Simulation & Visualization

# ErLam Toolkit: The Language

```
<Expression> ::= <Variable>
                | <Integer>
                | 'newchan'
                | '(' <Expression> ')',
                | <Expression> <Expression>
                | 'if' <Expression> <Expression> <Expression>
                | 'swap' <Expression> <Expression>
                | 'spawn' <Expression>
                | 'fun' <Variable> '.' <Expression>
```

# ErLam Toolkit: The Language

```
elib
  // ...
  ignore = (fun _.(fun y.y));
  omega = (fun x.(x x));
  // ...
  add = _erl[2]{ fun(X) when is_integer(X) ->
                  fun(Y) when is_integer(Y) ->
                      X+Y
                  end
                end
            };
  // ...
bile
```

# ErLam Toolkit: The Language

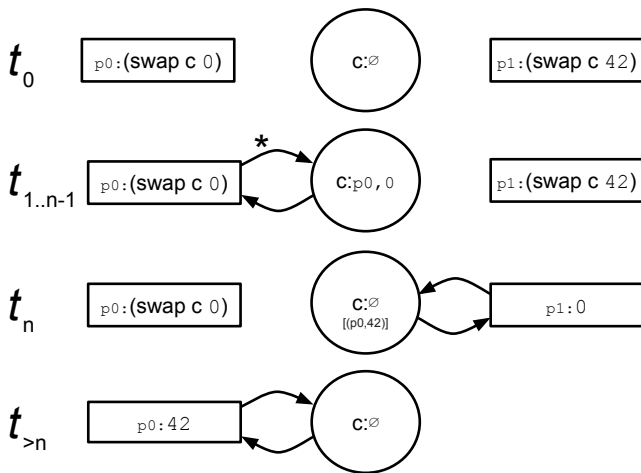
Example Application: Simple Swap

```
(fun c.  
  (ignore  
    (spawn (fun _.(swap c 42)))  
    (swap c 0))  
newchan)
```



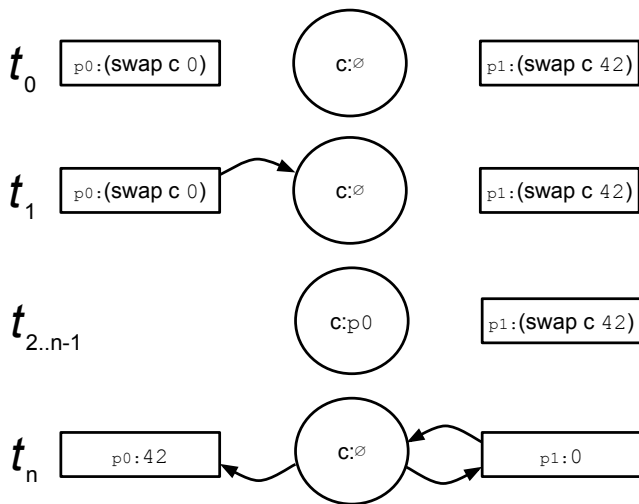
# ErLam Toolkit: Channel Implementations

## Process Blocking Swap



# ErLam Toolkit: Channel Implementations

## Process Absorption Swap



# ErLam Toolkit: Simulation & Visualization

## System Behaviours:

- Degree of Parallelism
- Consistency of Cooperation
- Degree of Longevity/Interactivity
- Partial System Cooperativity

## Logging & Report Generation



# ErLam Toolkit: Simulation & Visualization

## System Behaviours: Degree of Parallelism

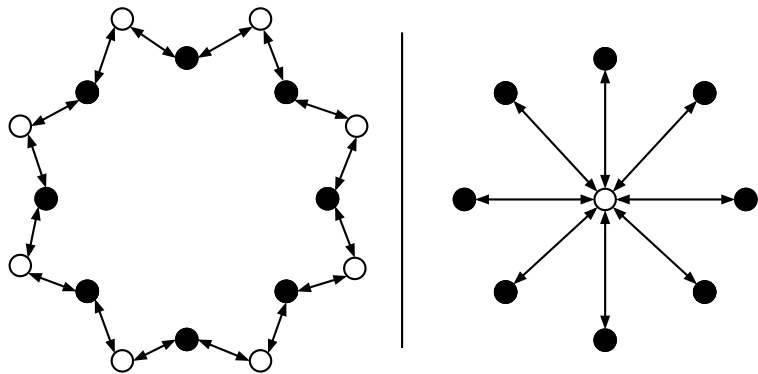


Figure:  $PRing_N$ , and  $ClusterComm_{(N,1)}$  primitives to test degree of parallelism.

# ErLam Toolkit: Simulation & Visualization

System Behaviours: Consistency of Cooperation

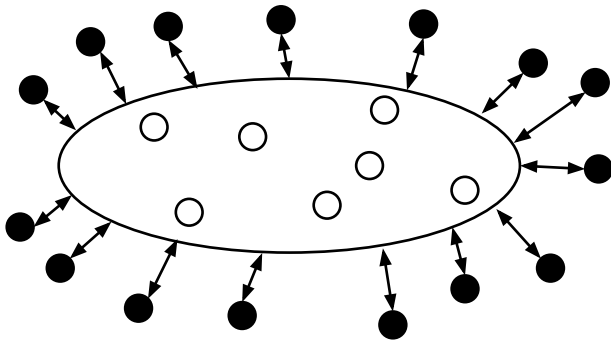


Figure:  $ClusterComm_{(N,M)}$  to test effect of consistency on scheduler.

# ErLam Toolkit: Simulation & Visualization

System Behaviours: Degree of Interactivity

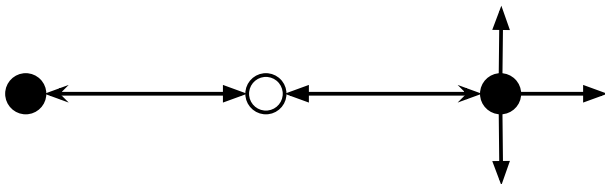


Figure:  $UserInput_{(T,C)}$ , simulates user interaction or a number ( $C$ ) of external/timed ( $T$ ) events.

# ErLam Toolkit: Simulation & Visualization

System Behaviours: Degree of Interactivity

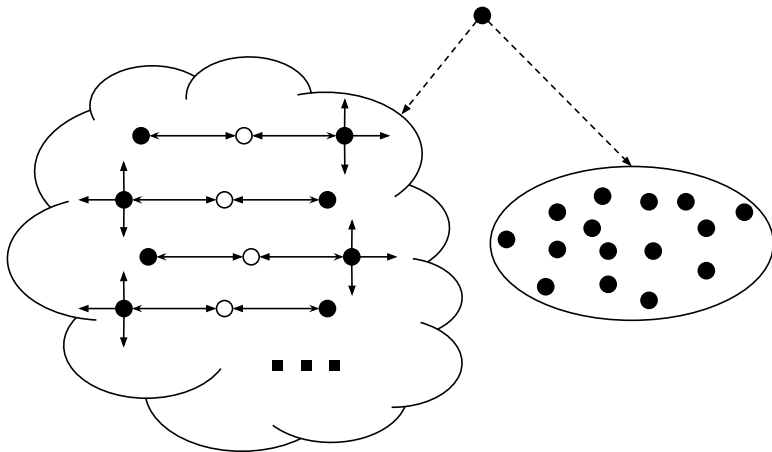


Figure:  $Interactivity_{(N,M)}$ , composure of  $ChugMachine_N$  (Cloud), and  $M$  instances of  $UserInput_{(5,2)}$ .

# ErLam Toolkit: Simulation & Visualization

System Behaviours: Partial System Cooperativity




Figure:  $PTree_{(W,N)}$ , a composure of  $W$   $ClusterComm_{(N,1)}$  instances running concurrently.

# ErLam Toolkit: Simulation & Visualization

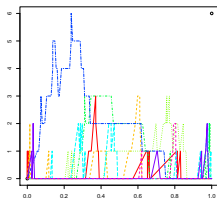
## Logging & Report Generation

Things we could log:

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts
- State of the Scheduler (waiting/running)
- Channel State (Blocked/Unblocked)
- ...

# ErLam Toolkit: Simulation & Visualization

**Queue Size**



**Reduc. Density**



**Channel State**



# Scheduler Implementations

## 3 Scheduler Implementations

- Example Schedulers
- Feedback Mechanisms



# Scheduler Implementations: Example Schedulers

- (MTRRGQ) - Round-Robin with Single Global Queue
  - All LPUs share a Process Queue.
- (MTRRWS-SQ) - Round-Robin with Work-Stealing via Direct Access
  - All LPUs have their own Process Queue.
  - LPUs can steal processes by grabbing them off the end of another LPU's queue.

# Scheduler Implementations: Feedback Mechanisms

Three types of mechanics:

- Longevity-Based Batching
- Channel Pinning
- Bipartite-Graph Aided Sorting

# Scheduler Implementations: Feedback Mechanisms

## Longevity-Based Batching

- Choose via Round-Robin
  - from batch rather than queue
  - keeps track of number of rounds (batch size)
- Work-Steal whole batches
- Spawn to batch unless:  $|b_i| \geq B$ 
  - Make singleton with new process.
  - Push parent and child into new batch.

# Scheduler Implementations: Feedback Mechanisms

## Longevity-Based Batching

- Choose via Round-Robin
  - from batch rather than queue
  - keeps track of number of rounds (batch size)
- Work-Steal whole batches
- Spawn to batch unless:  $|b_i| \geq B$ 
  - Make singleton with new process.
  - Push parent and child into new batch.

GOAL: Can batching based on longevity account for fine/coarse parallelism in application?

# Scheduler Implementations: Feedback Mechanisms

## Channel-Pinning

- Upon call to *newchan*, pin to LPU based on spread algorithm:
  - *same* - LPU *newchan* is called is where it is pinned.
  - *even* - Cycle through LPUs and pin based on that.
  - ...
- Work-steal based on channel that's been pinned to you.

# Scheduler Implementations: Feedback Mechanisms

## Channel-Pinning

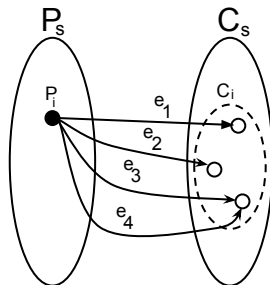
- Upon call to *newchan*, pin to LPU based on spread algorithm:
  - *same* - LPU *newchan* is called is where it is pinned.
  - *even* - Cycle through LPUs and pin based on that.
  - ...
- Work-steal based on channel that's been pinned to you.

GOAL: Can an *even*-like spread increase early saturation?

# Scheduler Implementations: Feedback Mechanisms

## Bipartite-Graph Aided Sorting


- Based on Round-Robin & Work-stealing
- Keep track of events which may effect cooperativity:
  - Spawning
  - Blocking/Unblocking
  - Steals
- If number of events over some threshold, re-sort.



# Scheduler Implementations: Feedback Mechanisms

## Bipartite-Graph Aided Sorting

- Based on Round-Robin & Work-stealing
- Keep track of events which may effect cooperativity:
  - Spawning
  - Blocking/Unblocking
  - Steals
- If number of events over some threshold, re-sort.



GOAL: Are alternate channel implementations worth exploration?



## | 4 Results

# Results:

## Longevity-Based Batching

- *Can batching based on longevity recognize fine/coarse parallelism in an application?*

## Channel Pinning

- *Can an even-like spread increase early saturation?*

## Bipartite-Graph Aided Sorting

- *Are alternate channel implementations worth exploration?*

# Results: Longevity-Based Batching

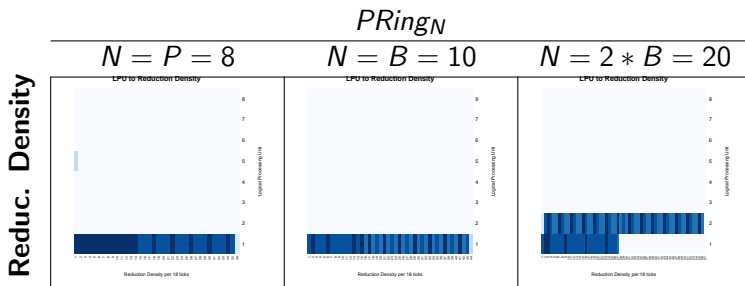

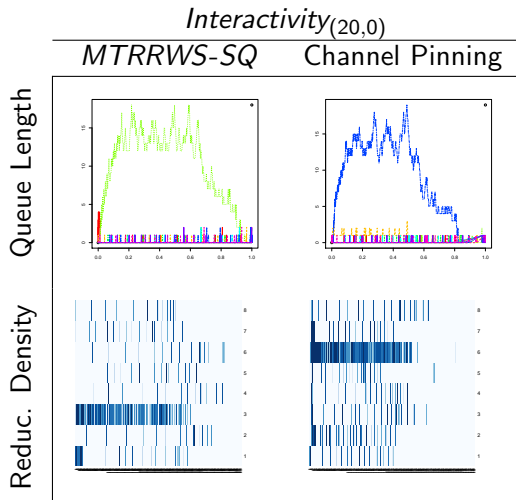


Table: Comparison of different sized  $PRing_N$  on the Longevity Batching Scheduler with batch size  $B = 10$ .

# Results: Longevity-Based Batching



# Results: Channel Pinning




# Results: Bipartite-Aided Graph Sorting

## Parallel Fibonacci


### *MTRRWS-SQ*

Channel State Over Time



### Sorting Scheduler

Channel State Over Time



# Conclusions & Future Work

## 5 Conclusions & Future Work

- ErLam Toolkit
- Cooperative Schedulers
- Cooperativity as a Metric

# Conclusions & Future Work: ErLam Toolkit

- Test Primitives were nicely composable process behaviours.
  - More research into generating behaviours.
  - More compositions: PTree with Rings.
- Log generation, lots of overhead, but good observations.



# Conclusions & Future Work: Cooperative Schedulers

## ■ Longevity Batching:

- Would benefit from heuristic based Quantum selection.
- As it stands, limited gain from longevity recognition.

## ■ Channel Pinning:

- Promising saturation and work-stealing mechanic.

## ■ Bipartite-Graph Aided Sorting:

- Supprising results on MapReduce style applications.
- Worth studying Blocking-Channels further for gains from sorting.

# Conclusions & Future Work: Cooperativity as a Metric

- Possible to recognize and benefit from.
- The three example mechanics are promising and can be extended for practice modern languages, despite simplistic simulation language.
- More to explore:
  - Alternate Message-Passing Types (Asymmetric?)
  - ...

# Questions/Comments?

# Questions/Comments?

Thank You!

# Links

- <https://github.com/dstar4138/erlam>
- [https://github.com/dstar4138/thesis\\_cooperativity](https://github.com/dstar4138/thesis_cooperativity)
- <http://dstar4138.com>