

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Alexander Dean

Rochester Institute of Technology
Golisano College of Computer and Information Sciences

August 12, 2014

Process Cooperativity as a Feedback Metric

in Concurrent Message-Passing Languages

1 Background

- Cooperativity
- Message Passing
- Runtime Scheduling

2 ErLam

- The Language
- Channel Implementations
- Simulation & Visualization

3 Scheduler Implementations

- Example Schedulers
- Feedback Schedulers

4 Results

5 Conclusions & Future Work

- ErLam Toolkit
- Cooperative Schedulers
- Cooperativity as a Metric



Current Section:

1 Background

- Cooperativity
- Message Passing
- Runtime Scheduling



Background: Cooperativity

What is Process Cooperativity?



Background: Cooperativity

What does Cooperativity give us?



Background: Message Passing

We use a Symmetric, Synchronous, Message-Passing Primitive:

swap

- Purely captures cooperation of processes by synchronizing on the shared channel.
- Ultimately can be extended to take into account:
 - Directionality
 - Asynchrony

Background: Runtime Scheduling

- Schedulers can be defined in a discrete manner:
 - 1 *Choose* a process from set
 - 2 *Reduce* it
 - 3 *Update* private scheduler state
- Statistics can be gathered at every step about process:
 - Number of channel partners,
 - Timestamp of last run,
 - Number of reductions, *etc.*
- *What statistics are necessary for recognizing cooperativity?*

2 ErLam

- The Language
- Channel Implementations
- Simulation & Visualization

ErLam: The Language

```
<Expression> ::= <Variable>
                | <Integer>
                | 'newchan'
                | '(' <Expression> ')',
                | <Expression> <Expression>
                | 'if' <Expression> <Expression> <Expression>
                | 'swap' <Expression> <Expression>
                | 'spawn' <Expression>
                | 'fun' <Variable> '.' <Expression>
```

ErLam: The Language

```
elib
  // ...
  omega = (fun x.(x x));
  // ...
  add = _erl[2]{ fun(X) when is_integer(X) ->
                  fun(Y) when is_integer(Y) ->
                      X+Y
                  end
                end
              };
  // ...
bile
```

ErLam: The Language


Example Application: Parallel Fibonacci

```
// pfib.els -  
fun N.  
  (omega fun f,m.(  
    if (leq m 1)  
      m  
      (merge fun _.(f f (sub m 1))  
              fun _.(f f (sub m 2))  
              add)) N)
```

\$ els pfib.els (Compile the script)


\$./pfib.ex -r 10 (Finds the 10th Fibonacci number)

ErLam: Channel Implementations



Process Blocking Swap

ErLam: Channel Implementations



Process Absorption Swap

ErLam: Simulation & Visualization

- System Behaviours:
 - Degree of Parallelism
 - Consistency of Cooperation
 - Degree of Longevity/Interactivity
 - Partial System Cooperativity
- Logging & Report Generation

ErLam: Simulation & Visualization

System Behaviours: Degree of Parallelism




Figure: $PRing_N$, and $ClusterComm_{(N,1)}$ primitives to test degree of parallelism.

ErLam: Simulation & Visualization

System Behaviours: Consistency of Cooperation




Figure: $ClusterComm_{(N,M)}$ to test effect of consistency on scheduler.

ErLam: Simulation & Visualization

System Behaviours: Degree of Interactivity




Figure: $UserInput_{(T,C)}$, simulates user interaction or a number (C) of external/timed (T) events.

ErLam: Simulation & Visualization

System Behaviours: Degree of Interactivity




Figure: $Interactivity_{(N,M)}$, composure of $ChugMachine_N$ (Cloud), and M instances of $UserInput_{(5,2)}$.

ErLam: Simulation & Visualization

System Behaviours: Partial System Cooperativity




Figure: $PTree_{(W,N)}$, a composure of W $ClusterComm_{(N,1)}$ instances running concurrently.

ErLam: Simulation & Visualization

Logging & Report Generation

Things we could log:

ErLam: Simulation & Visualization

Logging & Report Generation

Things we could log:

- Process Queue Size (per LPU)

ErLam: Simulation & Visualization

Logging & Report Generation

Things we could log:

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts

ErLam: Simulation & Visualization

Logging & Report Generation

Things we could log:

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts
- State of the Scheduler (waiting/running)

ErLam: Simulation & Visualization

Logging & Report Generation

Things we could log:

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts
- State of the Scheduler (waiting/running)
- Channel State (Blocked/Unblocked)

ErLam: Simulation & Visualization

Logging & Report Generation

Things we could log:

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts
- State of the Scheduler (waiting/running)
- Channel State (Blocked/Unblocked)
- ...

Current Section:

3 Scheduler Implementations

- Example Schedulers
- Feedback Schedulers

Scheduler Implementations: Example Schedulers

- Single-Thread Dual-Queue
 - Translation of CML scheduler
- Round-Robin with Single Global Queue (MTRRGQ)
 - No Work-Stealing
 - # LPUs can vary
- Round-Robin with Work-Stealing (MTRRWS)
 - Steal via direct access
 - Steal via advertisement

Scheduler Implementations: Feedback Schedulers

Three types of mechanics:

- Longevity-Based Batching
- Channel Pinning
- Bipartite-Graph Aided Sorting

Scheduler Implementations: Feedback Schedulers

Longevity-Based Batching

- Choose via Round-Robin
 - from batch rather than queue
 - keeps track of number of rounds (batch size)
- Work-Steal whole batches
- Spawn to batch unless: $|b_i| \geq B$
 - Make singleton with new process.
 - Push parent and child into new batch.

Scheduler Implementations: Feedback Schedulers

Longevity-Based Batching

- Choose via Round-Robin
 - from batch rather than queue
 - keeps track of number of rounds (batch size)
- Work-Steal whole batches
- Spawn to batch unless: $|b_i| \geq B$
 - Make singleton with new process.
 - Push parent and child into new batch.

GOAL: Can batching based on longevity account for fine/coarse parallelism in application?

Scheduler Implementations: Feedback Schedulers

Channel-Pinning

- Upon call to *newchan*, pin to LPU based on spread algorithm:
 - *same* - LPU *newchan* is called is where it is pinned.
 - *even* - Cycle through LPUs and pin based on that.
 - ...
- Work-steal based on channel that's been pinned to you.

Scheduler Implementations: Feedback Schedulers

Channel-Pinning


- Upon call to *newchan*, pin to LPU based on spread algorithm:
 - *same* - LPU *newchan* is called is where it is pinned.
 - *even* - Cycle through LPUs and pin based on that.
 - ...
- Work-steal based on channel that's been pinned to you.

GOAL: Can an *even*-like spread increase early saturation?

Scheduler Implementations: Feedback Schedulers

Bipartite-Graph Aided Sorting


- Based on Round-Robin & Work-stealing
- Keep track of events which may effect cooperativity:
 - Spawning
 - Blocking/Unblocking
 - Steals
- If number of events over some threshold, re-sort.



Scheduler Implementations: Feedback Schedulers

Bipartite-Graph Aided Sorting

- Based on Round-Robin & Work-stealing
- Keep track of events which may effect cooperativity:
 - Spawning
 - Blocking/Unblocking
 - Steals
- If number of events over some threshold, re-sort.



GOAL: Are alternate channel implementations worth exploration?

| 4 Results

Results:

Longevity-Based Batching

- *Can batching based on longevity recognize fine/coarse parallelism in an application?*

Channel Pinning

- *Can an even-like spread increase early saturation?*

Bipartite-Graph Aided Sorting

- *Are alternate channel implementations worth exploration?*

Results: Longevity-Based Batching


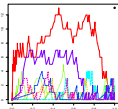
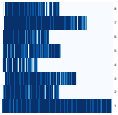
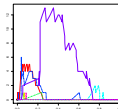

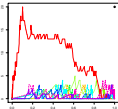
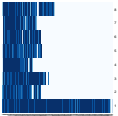
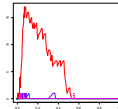

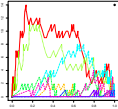
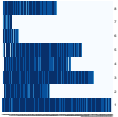
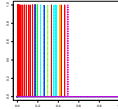
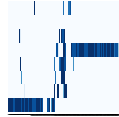
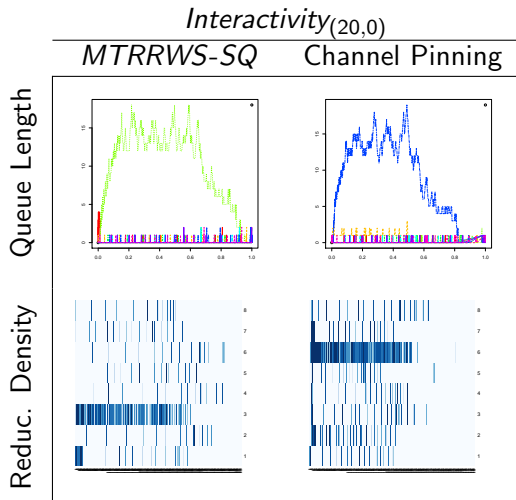


Table: Comparison of different sized $PRing_N$ on the Longevity Batching Scheduler with batch size $B = 10$.

Results: Longevity-Based Batching

Time Quantum	$P\text{Tree}_{(4,8)}$		Long. Batcher	
	Chan. State	Reduc. Density	Chan. State	Reduc. Density
$R = 30$				
$R = 40$				
$R = 50$				

Results: Channel Pinning




Results: Bipartite-Aided Graph Sorting

Parallel Fibonacci


MTRRWS-SQ

Channel State Over Time



Sorting Scheduler

Channel State Over Time



5 Conclusions & Future Work

- ErLam Toolkit
- Cooperative Schedulers
- Cooperativity as a Metric

Conclusions & Future Work: ErLam Toolkit

- Test Primitives were nicely composable process behaviours.
 - More research into generating behaviours.
 - More compositions: PTree with Rings.
- Log generation, lots of overhead, but good observations.

Conclusions & Future Work: Cooperative Schedulers

■ Longevity Batching:

- Would benefit from heuristic based Quantum selection.
- As it stands, limited gain from longevity recognition.

■ Channel Pinning:

- Promising saturation and work-stealing mechanic.

■ Bipartite-Graph Aided Sorting:

- Supprising results on MapReduce style applications.
- Worth studying Blocking-Channels further for gains from sorting.

Conclusions & Future Work: Cooperativity as a Metric

- Possible to recognize and benefit from.
- The three example mechanics are promising and can be extended for practice modern languages, despite simplistic simulation language.
- More to explore:
 - Alternate Message-Passing Types (Asymmetric?)
 - ...

Questions/Comments?

Questions/Comments?

Thank You!

Links

- <https://github.com/dstar4138/erlam>
- https://github.com/dstar4138/thesis_cooperativity
- <http://dstar4138.com>