

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

August 12, 2014

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

- └ Process Cooperativity as a Feedback Metric

Process Cooperativity as a Feedback Metric

in Concurrent Message-Passing Languages

- 1 Background
 - Runtime Scheduling
 - Cooperativity
 - Message Passing
- 2 Erlang Toolkit
 - The Language
 - Channel Implementations
 - Simulation & Visualization
- 3 Scheduler Implementations
 - Example Schedulers
 - Feedback Mechanisms
- 4 Results
- 5 Conclusions & Future Work
 - Erlang Toolkit
 - Cooperative Schedulers
 - Cooperativity as a Metric

- 1 Background
 - Runtime Scheduling
 - Cooperativity
 - Message Passing
 - 2 ErLam Toolkit
 - The Language
 - Channel Implementations
 - Simulation & Visualization
 - 3 Scheduler Implementations
 - Example Schedulers
 - Feedback Mechanisms
 - 4 Results
 - 5 Conclusions & Future Work
 - ErLam Toolkit
 - Cooperative Schedulers
 - Cooperativity as a Metric

- Mouthful of a title, so I'll break it up:
 1. Runtime scheduling, to give some grounding in the area of study.
 2. Cooperativity, what it is and motivation to use it.
 3. Message Passing, because, as it turns out, it's a nice abstraction for our purpose of capturing cooperativity.
- The core of the work revolves around the toolkit I built.
 - A language/compiler/runtime/testing-framework
 - But also a *Simulator* which has a plug-and-play scheduler API. It let me test schedulers on a common test bed.
- Next, go over the list of schedulers & feedback mechanisms.
- Results, Conclusions, & Future Work.

Background

2014-08-12

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Background
 - Background

2014-08-12

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Background
 - Background

2014-08-12

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Background
 - Background

2014-08-12

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Background
 - Background

2014-08-12

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Background
 - Background

- Background
 - Runtime Scheduling
 - Cooperativity
 - Message Passing

- Background
 - Runtime Scheduling
 - Cooperativity
 - Message Passing

Introduce the new section:

- 1 Background
 - Runtime Scheduling
 - Cooperativity
 - Message Passing

- Introduce the new section:
- 1 Background
 - Runtime Scheduling
 - Cooperativity
 - Message Passing

Background: Runtime Scheduling

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

- Background

-Runtime Scheduling

-Background: Runtime Scheduling

- Schedulers can be defined in a discrete manner:
 - 1 Choose a process from set,
 - 2 Reduce it,
 - 3 Update private scheduler state.

- Schedulers can be defined in a discrete manner:

- 1 Choose a process from set,
- 2 Reduce it,
- 3 Update private scheduler state.

- We can look at process schedulers like a function:
 - Takes a set of processes, and some private state.
 - Job of the function is to choose a process, and run it for a bit.
 - Then, based on what happened while running process, we update the state.
- Big questions: How are we choosing a process? What should effect our decision?

2014-08-12

Background

- Background: Runtime Scheduling

- Schedulers can be defined in a discrete manner:
 - 1 Choose a process from set,
 - 2 Reduce it,
 - 3 Update private scheduler state.
- Statistics can be gathered at every step about process:
 - Timestamp of last run,
 - Number of reductions, etc.

- We can look at process schedulers like a function:
 - Takes a set of processes, and some private state.
 - Job of the function is to choose a process, and run it for a bit.
 - Then, based on what happened while running process, we update the state.
- Big questions: How are we choosing a process? What should effect our decision?
- Timestamp of last run? →
 - Choose always most recent, it's a batch scheduler.
 - Choose oldest, we get something called Round-Robin.
- Number of reductions? → longevity = might want to give someone else a go.
- What are useful, and what do they tell us about the state of the system? Well this leads us to cooperativity.

2014-08-12

- Background

-Background: Runtime Scheduling

- Schedulers can be defined in a discrete manner:
 - 1 Choose a process from set,
 - 2 Reduce it,
 - 3 Update private scheduler state.
- Statistics can be gathered at every step about process:
 - Timestamp of last run,
 - Number of reductions, etc.
- What statistics are useful?

- We can look at process schedulers like a function:
 - Takes a set of processes, and some private state.
 - Job of the function is to choose a process, and run it for a bit.
 - Then, based on what happened while running process, we update the state.
- Big questions: How are we choosing a process? What should effect our decision?
- Timestamp of last run? →
 - Choose always most recent, it's a batch scheduler.
 - Choose oldest, we get something called Round-Robin.
- Number of reductions? → longevity = might want to give someone else a go.
- What are useful, and what do they tell us about the state of the system? Well this leads us to cooperativity.

Background: Cooperativity

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

└ Background

└ Cooperativity



└ Background: Cooperativity

Background: Cooperativity

What is Process Cooperativity?



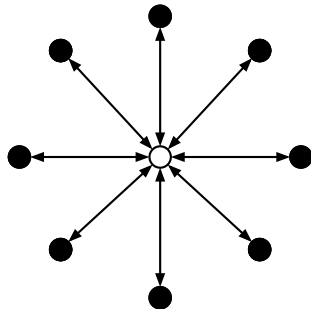

What is Process Cooperativity?



- What is Process Cooperativity?
- White = channel & black = a process.
- Can think of channel a mechanism for passing information between processes.
 - These are nice functional abstractions of things like locks, shared-memory, *etc.*
- Left: Cloud of processes with no interaction.
- Right: We see a definite structure caused by some sharing of information. This is the core of recognizing cooperation, namely, recognizing these structures when they exist.

Background: Cooperativity

What does Cooperativity give us?

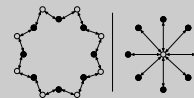


2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

- Background
 - Cooperativity
 - Background: Cooperativity

Background: Cooperativity



What does Cooperativity give us? What's the difference in the behaviour of cooperation in the left/right applications?

- Left: A Ring,
 - the level of parallelism is nearly nil.
 - Each process is cooperating yes, but granularity is very fine.
- Right: A Star,
 - the level of parallelism is nearly full.
 - Each process is cooperating, **not reliant on more than one** other process.
- In both, the whole system is communicating, but with cooperation, we can find the level of parallelism possible.

Next: Knowing this, how can we recognize cooperativity? Seems to be all about recording interactions with the channel.

Background: Message Passing

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

- Background

- Message Passing

-Background: Message Passing

We use a Symmetric, Synchronous, Message-Passing Primitive

swap

- Purely captures cooperation of processes through synchronizing on a shared channel.

We use a Symmetric, Synchronous, Message-Passing Primitive:

swap

- Purely captures cooperation of processes through synchronizing on a shared channel.

- Symmetric, Synchronous, Message-Passing primitive.
- Symmetric:
 - Only one message passing primitive: SWAP
- Synchronous:
 - Blocks until it's partner gets there.
- Purely captures cooperation: Simple synchronization representation.
- This is really what I based the language on.
- So, what does the rest of the language look like.

2 ErLam Toolkit

- The Language
- Channel Implementations
- Simulation & Visualization

Introduce the new section:

- 2 ErLam Toolkit
 - The Language
 - Channel Implementations
 - Simulation & Visualization

ErLam Toolkit: The Language

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

- ErLam Toolkit

-The Language

-ErLam Toolkit: The Language

```

<Expression> ::= <Variable>
               | <Integer>
               | 'neuchan'
               | '(' <Expression> ')'
               | <Expression> <Expression>
               | 'if' <Expression> <Expression> <Expression>
               | 'map' <Expression> <Expression>
               | 'apmap' <Expression>
               | 'fun' <Variable> '.' <Expression>

```

```
<Expression> ::= <Variable>
                | <Integer>
                | 'newchan'
                | '(' <Expression> ')',
                | <Expression> <Expression>
                | 'if' <Expression> <Expression> <Expression>
                | 'swap' <Expression> <Expression>
                | 'spawn' <Expression>
                | 'fun' <Variable> '.' <Expression>
```

- Extremely simple on purpose (5 keywords).
- Issue now began to be how to build up primitive test behaviours.
- Made a library which allowed for built ins.

ErLam Toolkit: The Language

```

elib
    // ...
    ignore = (fun _.(fun y.y));
    omega = (fun x.(x x));
    // ...
    add = _erl[2]{ fun(X) when is_integer(X) ->
                    fun(Y) when is_integer(Y) ->
                        X+Y
                    end
                end
            };
    // ...
bfile

```

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

- ErLam Toolkit

- The Language

-ErLam Toolkit: The Language

```

elim
// ...
ignore = (fun _.(fun y.y));
omega = (fun x.(x x));
// ...
add = _erl[2]{ fun(X) when is_integer(X) ->
                                fun(Y) when is_integer(Y) ->
                                    X+Y
                                end
                            };
// ...
bile

```

- There's options for built-ins as well as macros.
- Built-ins are raw Erlang, gets wrapped up into AST, and still "reduces" the same (*i.e.* no multi-variable functions).

ErLam Toolkit: The Language

Example Application: Simple Swap

```
(fun c.  
  (ignore  
    (spawn (fun _.(swap c 42)))  
    (swap c 0))  
  newchan)
```



2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages


ErLam Toolkit

The Language

ErLam Toolkit: The Language

ErLam Toolkit: The Language
Example Application: Simple Swap

```
(fun c.  
  (ignore  
    (spawn (fun _.(swap c 42)))  
    (swap c 0))  
  newchan)
```



- Here is a simple application which:
 - Spawns a process to swap the number 42
 - Calls swap to get the value from the other process.
- Build up from here to simulate more complex behaviour.
- We can "do some work" before swapping, *etc.*
- This is how we built up or primitive test behaviours.

ErLam Toolkit: Channel Implementations

Process Blocking Swap

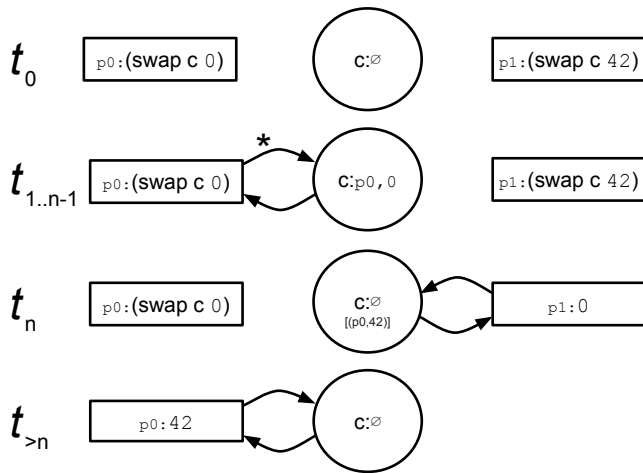
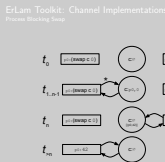
2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

ErLam Toolkit

Channel Implementations

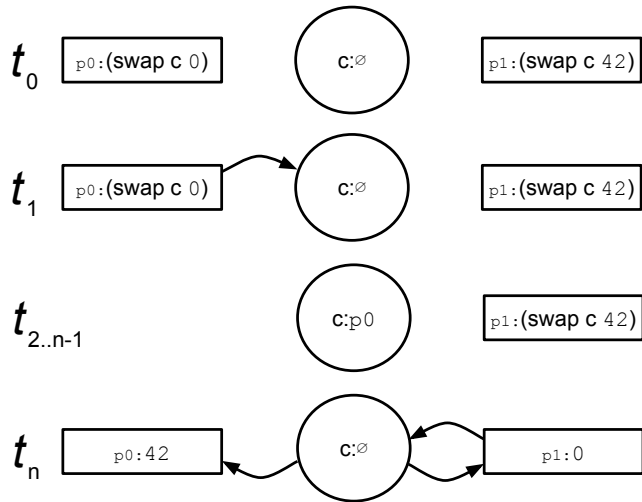
ErLam Toolkit: Channel Implementations



- Blocking: Maintains state of current and previous swap value until swap is completed.
- Mention expected effects on scheduler.
 - Scheduler will keep hold of the process, needs to recheck if blocked.
 - If all processes are communicating, large process queue of blocked processes.

ErLam Toolkit: Channel Implementations

Process Absorption Swap



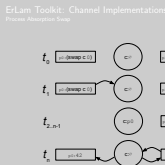
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12

ErLam Toolkit

Channel Implementations

ErLam Toolkit: Channel Implementations



1. Whole process gets absorbed by channel, away from scheduler.
 2. When the second process completes the swap, the process gets removed from channel.
 - The p_0 process can go back to its original scheduler,
 - OR to scheduler which unblocked it.
- Effects on scheduler:
 - Can loose or gain a extra process during communication.

ErLam Toolkit: Simulation & Visualization

- 2014-08-12 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - ErLam Toolkit
 - Simulation & Visualization
 - ErLam Toolkit: Simulation & Visualization

- 2014-08-12 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - ErLam Toolkit
 - Simulation & Visualization
 - ErLam Toolkit: Simulation & Visualization

- 2014-08-12 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - ErLam Toolkit
 - Simulation & Visualization
 - ErLam Toolkit: Simulation & Visualization

- 2014-08-12 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - ErLam Toolkit
 - Simulation & Visualization
 - ErLam Toolkit: Simulation & Visualization

- 2014-08-12 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - ErLam Toolkit
 - Simulation & Visualization
 - ErLam Toolkit: Simulation & Visualization

ErLam Toolkit: Simulation & Visualization

Primitive Testing Behaviours:

- Degree of Parallelism
- Partial System Cooperativity

Logging & Report Generation

ErLam Toolkit: Simulation & Visualization

Primitive Testing Behaviours:

- Degree of Parallelism
- Partial System Cooperativity

Logging & Report Generation

- ## ErLam Toolkit: Simulation & Visualization
- Primitive Testing Behaviours:
- Degree of Parallelism
 - Partial System Cooperativity
- Logging & Report Generation

ErLam Toolkit: Simulation & Visualization

Primitive Testing Behaviours:

- Degree of Parallelism
- Partial System Cooperativity

Logging & Report Generation

Primitive Testing Behaviours:

- Degree of Parallelism
- Partial System Cooperativity

Logging & Report Generation

- ## Primitive Testing Behaviours:
- Degree of Parallelism
 - Partial System Cooperativity
- ## Logging & Report Generation

Primitive Testing Behaviours:

- Degree of Parallelism
- Partial System Cooperativity

Logging & Report Generation

- To test cooperativity, we need a set of test behaviours.
- We want to analyze a large number of structure types.
- Two of them that I studied are:
 - Parallelism: Gets back to Ring vs Star, compare the two.
 - Full vs Partial Cooperation: Multiple groups of Stars or Rings.
- Finally, quickly, go over the current report generation that ErLam performs.

ErLam Toolkit: Simulation & Visualization

System Behaviours: Degree of Parallelism

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

ErLam Toolkit

Simulation & Visualization

ErLam Toolkit: Simulation & Visualization

ErLam Toolkit: Simulation & Visualization

System Behaviours: Degree of Parallelism

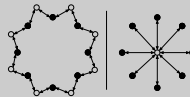


Figure: $PRing_N$ and $ClusterComm_{(N,1)}$ primitives to test degree of parallelism.

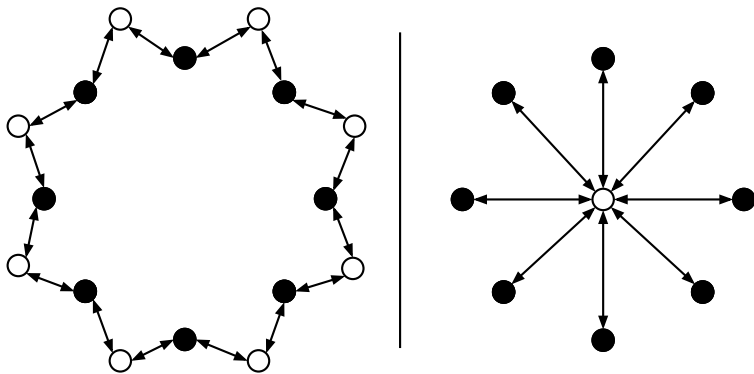


Figure: $PRing_N$ and $ClusterComm_{(N,1)}$ primitives to test degree of parallelism.

- Here are our Ring and Star behaviours, we call them by different names.
- We call the left, $PRing$, with the parameter N = number of processes.
- We call the right, $ClusterComm$ with two parameters, N like $PRing$, $M = 1$ in this case.
- Note the generalization of the Star is a set of M channels that are randomly accessed.

ErLam Toolkit: Simulation & Visualization

System Behaviours: Consistency of Cooperation

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

ErLam Toolkit

Simulation & Visualization

ErLam Toolkit: Simulation & Visualization

ErLam Toolkit: Simulation & Visualization
System Behaviours: Consistency of Cooperation

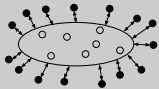


Figure: $ClusterComm_{(N,M)}$ to test effect of consistency on scheduler.

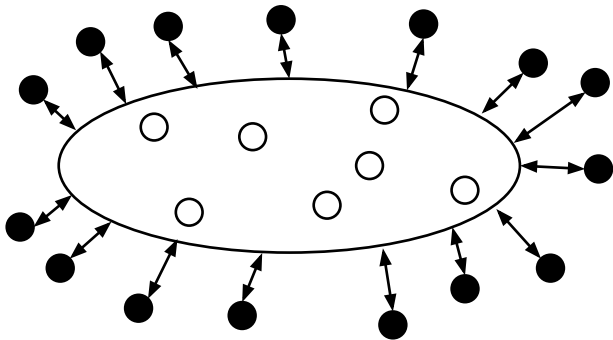


Figure: $ClusterComm_{(N,M)}$ to test effect of consistency on scheduler.

- Worst case scenario for C-C schedulers.
- However, this is still full system cooperativity. We want some minor work groups.

ErLam Toolkit: Simulation & Visualization

System Behaviours: Partial System Cooperativity

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

ErLam Toolkit

Simulation & Visualization

ErLam Toolkit: Simulation & Visualization

ErLam Toolkit: Simulation & Visualization
System Behaviours: Partial System Cooperativity

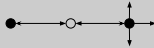


Figure: $UserInput_{(T,C)}$ simulates user interaction for a number (C) of external/timed (T) events.

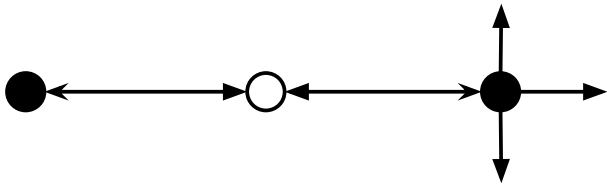


Figure: $UserInput_{(T,C)}$, simulates user interaction for a number (C) of external/timed (T) events.

- So we can make tiny work groups of two processes, like our simple swapping program earlier.
- The difference here is minor, we call this test `UserInput` because one process can hang for a set amount of time before triggering an event (effectively simulating user interaction).
- But we need multiple of these work groups, running concurrently.
- Run a set of them in parallel, and we could even inject a bit of overhead at will by running any number of extra processes.

ErLam Toolkit: Simulation & Visualization

System Behaviours: Partial System Cooperativity

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

ErLam Toolkit

Simulation & Visualization

ErLam Toolkit: Simulation & Visualization





Figure: $Interactivity_{(N,M)}$, composure of $ChugMachine_N$ (Cloud), and M instances of $UserInput_{(5,2)}$.

- So we can make tiny work groups of two processes, like our simple swapping program earlier.
- The difference here is minor, we call this test `UserInput` because one process can hang for a set amount of time before triggering an event (effectively simulating user interaction).
- But we need multiple of these work groups, running concurrently.
- Run a set of them in parallel, and we could even inject a bit of overhead at will by running any number of extra processes.

ErLam Toolkit: Simulation & Visualization

System Behaviours: Partial System Cooperativity

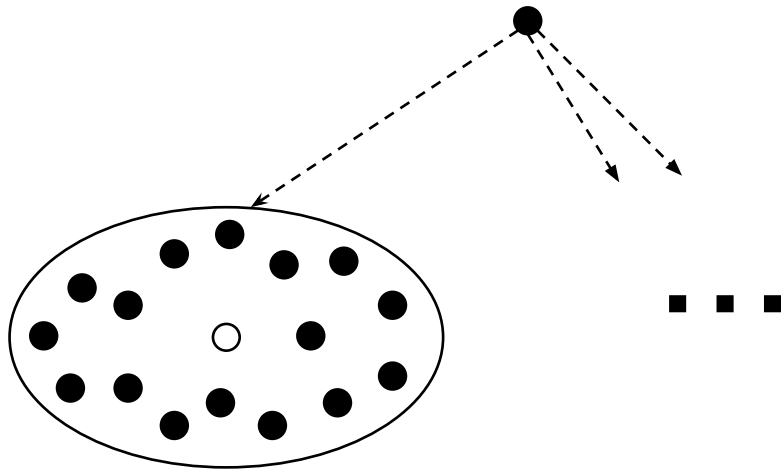


Figure: $PTree_{(W,N)}$, a composure of W $ClusterComm_{(N,1)}$ instances running concurrently.

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

ErLam Toolkit

Simulation & Visualization

ErLam Toolkit: Simulation & Visualization

ErLam Toolkit: Simulation & Visualization

System Behaviours: Partial System Cooperativity




Figure: $PTree_{(W,N)}$, a composure of W $ClusterComm_{(N,1)}$ instances running concurrently.

- But we can run ClusterComm or PRing in parallel.

Logging & Report Generation

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

-Simulation & Visualization

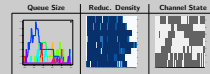
—ErLam Toolkit: Simulation & Visualization

Things we could log:

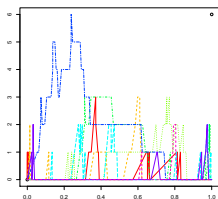
- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts
- State of the Scheduler (waiting/running)
- Channel State (Blocked/Unblocked)
- ...

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts
- State of the Scheduler (waiting/running)
- Channel State (Blocked/Unblocked)
- ...

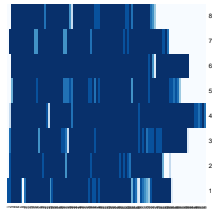
- Queue-Length: work-stealing mechanics and saturation ability.
- Tick-Action: Visualize the density of computation/communication.
- Sched-State: Useful for comparing stealing/process selection mechanics.
- Chan-State: Tracking interactivity, speed of unblock=attentive to cooperation.
- Of course there are more, but we limited ourselves to the above for initial testing purposes.



Queue Size



Reduc. Density



Channel State



- Three types of graphs:
 - Queue Size: X-axis is time, Y-axis is size of queue
 - Density charts: Darkness of the line represents fraction of ticks event happened in.
 - Channel State: dark=blocked, light=unblocked.

Scheduler Implementations

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Scheduler Implementations
 - Scheduler Implementations

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Scheduler Implementations
 - Scheduler Implementations

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Scheduler Implementations
 - Scheduler Implementations

Scheduler Implementations

- Scheduler Implementations
- Example Schedulers
- Feedback Mechanisms

Scheduler Implementations

- Scheduler Implementations
- Example Schedulers
- Feedback Mechanisms

Scheduler Implementations

- Scheduler Implementations
- Example Schedulers
- Feedback Mechanisms

3 Scheduler Implementations

- Example Schedulers
- Feedback Mechanisms

- # 3 Scheduler Implementations

 - Example Schedulers
 - Feedback Mechanisms

Introduce the new section:

- 3 Scheduler Implementations
 - Example Schedulers
 - Feedback Mechanisms

Introduce the new section:

- 3 Scheduler Implementations
 - Example Schedulers
 - Feedback Mechanisms

- Introduce the new section:
- 3 Scheduler Implementations
 - Example Schedulers
 - Feedback Mechanisms

2014-08-12

Scheduler Implementations

-Example Schedulers

Scheduler Implementations: Example Schedulers

- (MTRRGQ) - Round-Robin with Single Global Queue
 - All LPU's share a Process Queue.
- (MTRRWS-SQ) - Round-Robin with Work-Stealing via Direct Access
 - All LPU's have their own Process Queue.
 - LPU's can steal processes by grabbing them off the end of another LPU's queue.

1. RR w/ Global Queue: all synchronization around a single shared queue
2. RR w/ Work-Stealing: each scheduler gets their own queue but, they now need to steal work from others.
 - Implemented multiple types of work stealing, but we'll limit talk to one type:
 - Stealing directly from another LPUs by accessing the end of their process queue.

Scheduler Implementations: Feedback Mechanisms

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Scheduler Implementations
 - Feedback Mechanisms
 - Scheduler Implementations: Feedback Mechanisms

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Scheduler Implementations
 - Feedback Mechanisms
 - Scheduler Implementations: Feedback Mechanisms

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Scheduler Implementations
 - Feedback Mechanisms
 - Scheduler Implementations: Feedback Mechanisms

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Scheduler Implementations
 - Feedback Mechanisms
 - Scheduler Implementations: Feedback Mechanisms

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Scheduler Implementations
 - Feedback Mechanisms
 - Scheduler Implementations: Feedback Mechanisms

Scheduler Implementations: Feedback Mechanisms

Three types of mechanisms:

- Longevity-Based Batching
- Channel Pinning
- Bipartite-Graph Aided Sorting

- ## Scheduler Implementations: Feedback Mechanisms
- Three types of mechanisms:
- Longevity-Based Batching
 - Channel Pinning
 - Bipartite-Graph Aided Sorting

Three types of mechanics:

- Longevity-Based Batching
- Channel Pinning
- Bipartite-Graph Aided Sorting

- ## Three types of mechanics:
- Longevity-Based Batching
 - Channel Pinning
 - Bipartite-Graph Aided Sorting

- Instead of a single cooperativity-conscious scheduler, we implemented three mechanics which take cooperativity into account on top of the basic schedulers.

Scheduler Implementations: Feedback Mechanisms

Longevity-Based Batching

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

-Scheduler Implementations

-Feedback Mechanisms

-Scheduler Implementations: Feedback Mechanisms

- Choose via Round-Robin
 - from batch rather than queue
 - keeps track of number of rounds (batch size)
- Work-Steal whole batches
- Spawn to batch unless: $|b_i| \geq B$
 - Make singleton with new process.
 - Push parent and child into new batch.

- Batching processes based on longevity.
 - Based on $\text{occam-}\Pi$.
 - if a process communicates frequently then it will be batched (absorption), singleton if very computation-bound.
- We are normal RR but with one extra layer.
- If batch is too big during spawns we can:
 - Make singleton, best if child is needed to start work right away. Map-Reduce.
 - Make push-back, parent can get another chance to spawn more children sooner.
- GOAL here is to analyze the effects of grouping the frequently communicating.

Scheduler Implementations: Feedback Mechanisms

Longevity-Based Batching

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Scheduler Implementations

- Feedback Mechanisms

-Scheduler Implementations: Feedback Mechanisms

- Choose via Round-Robin
 - from batch rather than queue
 - keeps track of number of rounds (batch size)
- Work-Steal whole batches
- Spawn to batch unless: $|b_i| \geq B$
 - Make singleton with new process.
 - Push parent and child into new batch.

GOAL: Can batching based on longevity account for fine/coarse parallelism in application?

- Choose via Round-Robin
 - from batch rather than queue
 - keeps track of number of rounds (batch size)
- Work-Steal whole batches
- Spawn to batch unless: $|b_i| \geq B$
 - Make singleton with new process.
 - Push parent and child into new batch.

GOAL: Can batching based on longevity account for fine/coarse parallelism in application?

- Batching processes based on longevity.
 - Based on $\text{occam-}\Pi$.
 - if a process communicates frequently then it will be batched (absorption), singleton if very computation-bound.
- We are normal RR but with one extra layer.
- If batch is too big during spawns we can:
 - Make singleton, best if child is needed to start work right away. Map-Reduce.
 - Make push-back, parent can get another chance to spawn more children sooner.
- GOAL here is to analyze the effects of grouping the frequently communicating.

Channel-Pinning

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Scheduler Implementations

- Feedback Mechanisms

Scheduler Implementations: Feedback Mechanisms

- Upon call to *newchan*, pin to LPU based on spread algorithm:
 - *same* - LPU *newchan* is called is where it is pinned.
 - *even* - Cycle through LPUs and pin based on that.
 - ...
- Work-steal based on channel that's been pinned to you.

- Pin channels to LPUs.
 - Pinning a channel means to set a process affinity to a LPU based on the channels it uses.
 - Work-Stealing works like Go-Fish.
- Absorption channels would work well here to relocate processes based on channel usage.
- GOAL: can we improve on work-stealing by being selective.

- Upon call to `newchan`, pin to LPU based on spread algorithm
 - `same` - LPU `newchan` is called is where it is pinned.
 - `even` - Cycle through LPUs and pin based on that.
 - ...
- Work-steal based on channel that's been pinned to you.

Scheduler Implementations: Feedback Mechanisms

Channel-Pinning

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Scheduler Implementations

– Feedback Mechanisms

Scheduler Implementations: Feedback Mechanisms

- Upon call to `newchan`, pin to LPU based on spread algorithm
 - `same` - LPU `newchan` is called is where it is pinned.
 - `even` - Cycle through LPUs and pin based on that.
 - ...
 - Work-steal based on channel that's been pinned to you.
- GOAL: Can an even-like spread increase early saturation?

- Upon call to *newchan*, pin to LPU based on spread algorithm:
 - *same* - LPU *newchan* is called is where it is pinned.
 - *even* - Cycle through LPUs and pin based on that.
 - ...
- Work-steal based on channel that's been pinned to you.

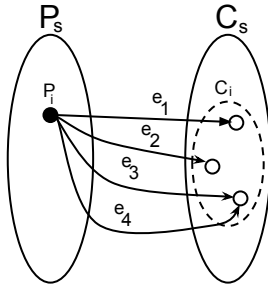
GOAL: Can an *even-like* spread increase early saturation?

- Pin channels to LPUs.
 - Pinning a channel means to set a process affinity to a LPU based on the channels it uses.
 - Work-Stealing works like Go-Fish.
- Absorption channels would work well here to relocate processes based on channel usage.
- GOAL: can we improve on work-stealing by being selective.

Scheduler Implementations: Feedback Mechanisms

Bipartite-Graph Aided Sorting

- Based on Round-Robin & Work-stealing
- Keep track of events which may effect cooperativity:
 - Spawning
 - Blocking/Unblocking
 - Steals
- If number of events over some threshold, re-sort.



2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

└ Scheduler Implementations

└ Feedback Mechanisms

└ Scheduler Implementations: Feedback Mechanisms

Scheduler Implementations: Feedback Mechanisms

Bipartite-Graph Aided Sorting

- Based on Round-Robin & Work-stealing
- Keep track of events which may effect cooperativity:
 - Spawning
 - Blocking/Unblocking
 - Steals
- If number of events over some threshold, re-sort.




- Based on RR and WS.
- Keep a list of all communications as a graph between set of processes and channels.
- Keep track of all queue events: Spawn, Block, Steal
- We can resort using some function with these edges.
- GOAL: If we block, we wont loose order. If we had absorption channels like the previous two schedulers, all sorting would be in vain.

Scheduler Implementations: Feedback Mechanisms

Bipartite-Graph Aided Sorting

- Based on Round-Robin & Work-stealing
- Keep track of events which may effect cooperativity:
 - Spawning
 - Blocking/Unblocking
 - Steals
- If number of events over some threshold, re-sort.



GOAL: Are alternate channel implementations worth exploration?

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

└ Scheduler Implementations

└ Feedback Mechanisms

└ Scheduler Implementations: Feedback Mechanisms

Scheduler Implementations: Feedback Mechanisms

Reactive Graph Aided Sorting

- Based on Round-Robin & Work-stealing
- Keep track of events which may effect cooperativity:
 - Spawning
 - Blocking/Unblocking
 - Steals
- If number of events over some threshold, re-sort.

GOAL: Are alternate channel implementations worth exploration?

- Based on RR and WS.
- Keep a list of all communications as a graph between set of processes and channels.
- Keep track of all queue events: Spawn, Block, Steal
- We can resort using some function with these edges.
- GOAL: If we block, we wont loose order. If we had absorption channels like the previous two schedulers, all sorting would be in vain.

Results

2014-08-12

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Results
 - Results

2014-08-12

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Results
 - Results

2014-08-12

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Results
 - Results

2014-08-12

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Results
 - Results

Results:

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Results
 - Results:

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Results
 - Results:

- 2014-08-12 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - └ Results
 - └ Results:

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Results
 - Results:

- Results
 - Longevity-Based Batching
 - Can batching based on longevity recognize fine/course parallelism in an application?
 - Channel Fitting
 - Can an even-like spread increase early saturation?
 - Bipartite-Graph Aided Sorting
 - Are alternate channel implementations worth exploration?

- Results
 - Longevity-Based Batching
 - Can batching based on longevity recognize fine/course parallelism in an application?
 - Channel Fitting
 - Can an even-like spread increase early saturation?
 - Bipartite-Graph Aided Sorting
 - Are alternate channel implementations worth exploration?

- Results
 - Longevity-Based Batching
 - Can batching based on longevity recognize fine/course parallelism in an application?
 - Channel Fitting
 - Can an even-like spread increase early saturation?
 - Bipartite-Graph Aided Sorting
 - Are alternate channel implementations worth exploration?

- Results
 - Longevity-Based Batching
 - Can batching based on longevity recognize fine/course parallelism in an application?
 - Channel Fitting
 - Can an even-like spread increase early saturation?
 - Bipartite-Graph Aided Sorting
 - Are alternate channel implementations worth exploration?

- Results
 - Longevity-Based Batching
 - Can batching based on longevity recognize fine/course parallelism in an application?
 - Channel Fitting
 - Can an even-like spread increase early saturation?
 - Bipartite-Graph Aided Sorting
 - Are alternate channel implementations worth exploration?

- Results
 - Longevity-Based Batching
 - Can batching based on longevity recognize fine/course parallelism in an application?
 - Channel Fitting
 - Can an even-like spread increase early saturation?
 - Bipartite-Graph Aided Sorting
 - Are alternate channel implementations worth exploration?

- Results
 - Longevity-Based Batching
 - Can batching based on longevity recognize fine/course parallelism in an application?
 - Channel Fitting
 - Can an even-like spread increase early saturation?
 - Bipartite-Graph Aided Sorting
 - Are alternate channel implementations worth exploration?

Longevity-Based Batching

- *Can batching based on longevity recognize fine/coarse parallelism in an application?*

- ## Longevity-Based Batching
- *Can batching based on longevity recognize fine/coarse parallelism in an application?*

Channel Pinning

- *Can an even-like spread increase early saturation?*

- ### Channel Pinning
- *Can an even-like spread increase early saturation?*

Bipartite-Graph Aided Sorting

- *Are alternate channel implementations worth exploration?*

- ## Bipartite-Graph Aided Sorting
- *Are alternate channel implementations worth exploration?*

- Remind about the goals of the talk.
- LBB: Would like to take advantage of the frequency of communication.

Results: Longevity-Based Batching


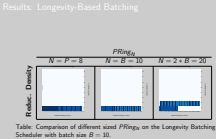


Table: Comparison of different sized $PRing_N$ on the Longevity Batching Scheduler with batch size $B = 10$.

- Early tests gave promising results.
- Here is $PRing_N$ which shows the reabsorption and containment on a single LPU as expected and hoped.
- So does batching based on longevity really recognize fine/coarse parallelism in an application?
- Sort of, if you know what the right time-quantum is to make that distinction.



Results:

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent
Message-Passing Languages

└ Results

└ Results:



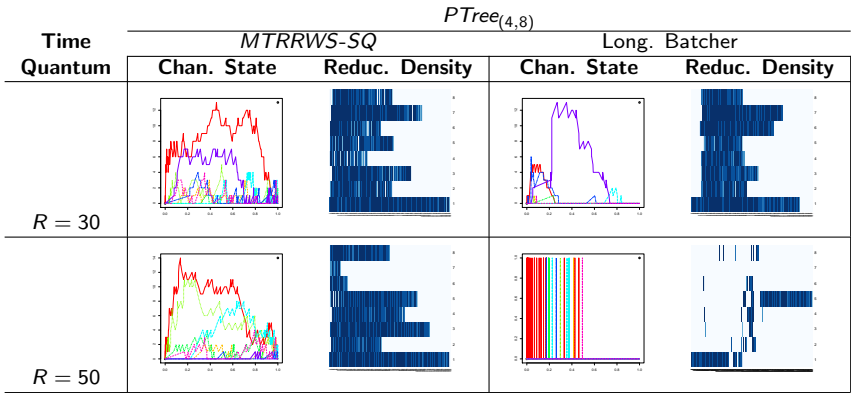

Reference Slide.

Results: Longevity-Based Batching

2014-08-12 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Results



Results: Longevity-Based Batching



- Comparison of $PTree_{(4,8)}$ running with the Longevity-Based Batching Scheduler and $MTRRWS-SQ$ at different time-quantums.
- Absorption channels help here to relocate processes.
- At lower time quantumms Long. Batcher starts to look like RRWS-SQ, however batching and absorption channels tend to lead to consolidation.
- At higher time quantumms Long. Batcher results in the originally expected work-groups. But it turns out to be inefficient due to lost chances of parallelism of each "star" of each group.
- Heuristical adjustment of the time-quantum would definitely be possible.
- NOTE: We don't capture overhead of stealing. Batching has obvious gains here.

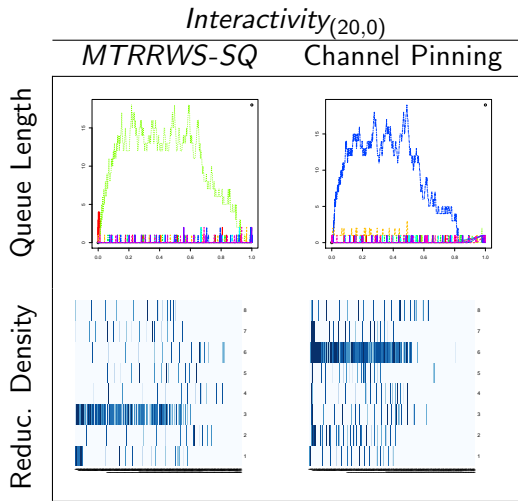
Results:

2014-08-12
Process Cooperativity as a Feedback Metric in Concurrent
Message-Passing Languages
└─ Results
└─ Results:



Reference Slide.

Results: Channel Pinning




2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Results

Results: Channel Pinning

Results: Channel Pinning



- Comparison of Uniform synchronization for *MTRRWS-SQ* and the Channel Pinning Scheduler on Absorption Channels.
- This used the *even* spread type.
- Note the speed at which it saturates all cores.
- Despite Naive WS, we still have decent spread.


Results:

2014-08-12


Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

└ Results

└ Results:



Reference Slide.



Results:




Results: Bipartite-Aided Graph Sorting

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Results

- Results:



- There is an issue with our next test.
- Our base primitive behaviours don't have a preferred order of execution.
- Ring does, but there is nothing parallel about it.

Results:

Results: Bipartite-Aided Graph Sorting


```
fun N.  
  (omega fun f,m.(  
    if (leq m 1)  
      m  
      (merge fun _.(f f (sub m 1))  
               fun _.(f f (sub m 2))  
               add)) N)
```

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Results

Results:



- We therefore used a parallel Fibonacci program.
- Merge function waits for each sub process to finish before running Add.
- If we sort based on who hasn't had a chance to communicate yet, we can preferentialize these MapReduce style applications, while not decrementing possible execution in parents.


Results: Bipartite-Aided Graph Sorting

2014-08-12

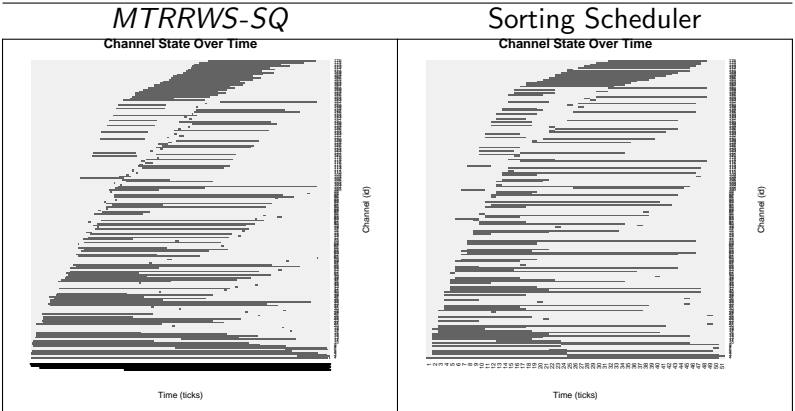
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Results

Results: Bipartite-Aided Graph Sorting



Parallel Fibonacci



- PFib has a strong reliance on order of execution.
- Channel State comparison of Parallel Fibonacci executed on *MTRRWS-SQ* and the Bipartite-Graph Aided Sorting Scheduler.
- Note the large reduction in number of ticks.

Conclusions & Future Work

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Conclusions & Future Work
 - Conclusions & Future Work

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Conclusions & Future Work
 - Conclusions & Future Work

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Conclusions & Future Work
 - Conclusions & Future Work

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Conclusions & Future Work
 - Conclusions & Future Work

Conclusions & Future Work

■ Erlang Toolkit

■ Cooperative Schedulers

■ Cooperativity as a Metric

■ Conclusions & Future Work

- Conclusions & Future Work
 - ErLam Toolkit
 - Cooperative Schedulers
 - Cooperativity as a Metric

- Conclusions & Future Work
 - ErLam Toolkit
 - Cooperative Schedulers
 - Cooperativity as a Metric

Introduce the new section:

- 5 Conclusions & Future Work
 - ErLam Toolkit
 - Cooperative Schedulers
 - Cooperativity as a Metric

Introduce the new section:

- 5 Conclusions & Future Work
 - ErLam Toolkit
 - Cooperative Schedulers
 - Cooperativity as a Metric

- Introduce the new section:
- 5 Conclusions & Future Work
 - ErLam Toolkit
 - Cooperative Schedulers
 - Cooperativity as a Metric

Conclusions & Future Work: ErLam Toolkit

- 2014-08-12 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - └─ Conclusions & Future Work
 - └─ ErLam Toolkit
 - └─ Conclusions & Future Work: ErLam Toolkit

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Conclusions & Future Work
 - ErLam Toolkit
 - Conclusions & Future Work: ErLam Toolkit

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Conclusions & Future Work
 - ErLam Toolkit
 - Conclusions & Future Work: ErLam Toolkit

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Conclusions & Future Work
 - ErLam Toolkit
 - Conclusions & Future Work: ErLam Toolkit

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Conclusions & Future Work
 - ErLam Toolkit
 - Conclusions & Future Work: ErLam Toolkit

- Test Primitives were nicely composable process behaviours.
 - More research into generating behaviours.
 - More compositions: FTree with Rings.
- Log generation, lots of overhead, but good observations.

- Test Primitives were nicely composable process behaviours.
 - More research into generating behaviours.
 - More compositions: PTree with Rings.
- Log generation, lots of overhead, but good observations.

- Overall pleased with simulator and achieved its goal.
- Future Work:
 - Generate more test primitives, a good library of them would be nice.
 - Compose them easier and more frequently. Perhaps generating work groups as PRing might have made a better comparison than Interactivity.
 - Clean up log generation (reduce overhead).
 - Process evaluation uses alpha-reduction (can be sped up substantially).
 - Make schedulers more adjustable (different spawn/yields/etc).
 - More Channel implementations

Conclusions & Future Work: Cooperative Schedulers

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Conclusions & Future Work

-Cooperative Schedulers

-Conclusions & Future Work: Cooperative Schedulers

- **Longevity Batching:**

- Would benefit from heuristic based Quantum selection.
- As it stands, limited gain from longevity recognition.

- **Channel Pinning:**

- Promising saturation and work-stealing mechanic

- **Bipartite-Graph Aided Sorting:**

- Surprising results on MapReduce style applications.
- Worth studying Blocking-Channels further for gains from sorting.

■ Longevity Batching:

- Would benefit from heuristic based Quantum selection.
- As it stands, limited gain from longevity recognition.

■ Channel Pinning:

- Promising saturation and work-stealing mechanic.

■ Bipartite-Graph Aided Sorting:

- Surprising results on MapReduce style applications.
- Worth studying Blocking-Channels further for gains from sorting.

- SS: Would benefit from merging with Channel Pinning. Increase likelihood that sorting puts channel partners close.
- LBB + CP might be interesting as channels could own batches.
- Overhead of Sorting? If implemented in a practical language, would it be worth it? Seems counter-intuitive but promising as is.

Conclusions & Future Work: Cooperativity as a Metric

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Conclusions & Future Work

-Cooperativity as a Metric

-Conclusions & Future Work: Cooperativity as a Metric

- Possible to recognize and benefit from.
- The three example mechanics are promising and can be extended for practice modern languages, despite simplistic simulation language.
- More to explore:
 - Alternate Message-Passing Types (Asymmetric?)
 - ...

- Promising scheduling mechanics
- More research is necessary in the message-passing implementation for it to be practical in common languages.
- (Asymmetric/Directionality would be first on the list)

- Possible to recognize and benefit from.
- The three example mechanics are promising and can be extended for practice modern languages, despite simplistic simulation language.
- More to explore:
 - Alternate Message-Passing Types (Asymmetric?)
 - ...

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent
Message-Passing Languages
└─ Questions

Questions/Comments?

Questions/Comments?

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent
Message-Passing Languages
└─ Questions

Questions/Comments?

Thank You!

Questions/Comments?

Thank You!

Links

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent
Message-Passing Languages

└─ Questions

└─ Links

Links

- <https://github.com/dstar4138/erlam>
- https://github.com/dstar4138/thesis_cooperativity
- <http://dstar4138.com>

- <https://github.com/dstar4138/erlam>
- https://github.com/dstar4138/thesis_cooperativity
- <http://dstar4138.com>