

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Alexander Dean

Rochester Institute of Technology
Golisano College of Computer and Information Sciences

August 12, 2014

Navigation icons: back, forward, search, etc.

2014-08-11 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
Alexander Dean
Rochester Institute of Technology
Golisano College of Computer and Information Sciences
August 12, 2014

- Thank Fluet, Heliotis, and Raj.
- Dedicate to parents, who are unable to be present.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-11 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

└ Process Cooperativity as a Feedback Metric

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
■ Background
■ Cooperativity
■ Message Passing
■ Runtime Scheduling
■ ErLam
■ The Language
■ Channel Implementations
■ Simulation & Visualization
■ Scheduler Implementations
■ Example Schedulers
■ Feedback Schedulers
■ Results
■ Conclusions & Future Work
■ ErLam Toolkit
■ Cooperative Schedulers
■ Cooperativity as a Metric

- 1 Background
 - Cooperativity
 - Message Passing
 - Runtime Scheduling
- 2 ErLam
 - The Language
 - Channel Implementations
 - Simulation & Visualization
- 3 Scheduler Implementations
 - Example Schedulers
 - Feedback Schedulers
- 4 Results
- 5 Conclusions & Future Work
 - ErLam Toolkit
 - Cooperative Schedulers
 - Cooperativity as a Metric

Navigation icons: back, forward, search, etc.

- Break apart title = Background.
- Then simulator and test primitives (process behaviours).
- Schedulers and how comparisons are made.
- Results, Conclusions (of both simulator and cooperativity).

Current Section:

2014-08-11 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

└ Background

└ Current Section:

Current Section
■ Background
■ Cooperativity
■ Message Passing
■ Runtime Scheduling


- 1 Background

- Cooperativity
- Message Passing
- Runtime Scheduling

Navigation icons: back, forward, search, etc.

Background: Cooperativity

What is Process Cooperativity?



Navigation icons: back, forward, search, etc.


2014-08-11
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
└ Background
└ Cooperativity
└ Background: Cooperativity

Background: Cooperativity
What is Process Cooperativity?

- Where white represents a channel and black represents a process.
- Channel = Source of synchronization (*i.e.* locks, abstractions, *etc.*)
- Left: Cloud of processes with no interaction.
- Right: Some sort of batch job? A bit of shared state amongst all processes these processes could all really be parallel or be competing, but in either case they cooperate.
- DEGREE OF COOPERATIVITY:
 - of *process*: flux of interaction with inter-proc comm method.
 - of *system*: rate of interaction between Ps and Cs as both flux in size.

Background: Cooperativity

What does Cooperativity give us?



Navigation icons: back, forward, search, etc.

2014-08-11
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
└ Background
└ Cooperativity
└ Background: Cooperativity

Background: Cooperativity
What does Cooperativity give us?

Whats the difference in cooperativity in the left and right set of processes?

- Left: A Ring, the level of parallelism is nearly nil. Each process is cooperating yes, but the granularity of the application is very fine.
- Right: A Star, the level of parallelism is nearly full. Each process is cooperating, and is **not reliant on more than one** other.

Overall, what can be gained by looking at cooperativity in terms of understanding the applications behaviour? Knowing the granularity of parallelism.

Next: what is a minimal inter-proc comm method?

Background: Message Passing

We use a Symmetric, Synchronous, Message-Passing Primitive:

swap

- Purely captures cooperation of processes by synchronizing on the shared channel.
- Ultimately can be extended to take into account:
 - Directionality
 - Asynchrony

Navigation icons: back, forward, search, etc.


2014-08-11
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
└ Background
└ Message Passing
└ Background: Message Passing

Background: Message Passing
We use a Symmetric, Synchronous, Message-Passing Primitive: swap
■ Purely captures cooperation of processes by synchronizing on the shared channel.
■ Ultimately can be extended to take into account:
■ Directionality
■ Asynchrony

- Async: while user code doesn't block, there is blocking in terms of the channel implementation. This is not indicated (in most cases) to the scheduler.
- Sync: The issue of process cooperation has been elevated to the process level for the scheduler to directly involve itself.
- Ultimately there is nothing stopping us from choosing the other types of message passing, but it would conflate the issue of process cooperativity if all we are after is whether two processes are cooperating on some task.

Next: How would we use coop as a feedback metric?
What does that mean?

Background: Runtime Scheduling

- # Background: Runtime Scheduling
- Schedulers can be defined in a discrete manner:
 - 1 Choose a process from set
 - 2 Reduce it
 - 3 Update private scheduler state
 - Statistics can be gathered at every step about process:
 - Number of channel partners,
 - Timestamp of last run,
 - Number of reductions, etc.
 - *What statistics are necessary for recognizing cooperativity?*
- 
- A set of small, faint navigation icons typically found in Beamer presentations, including symbols for back, forward, search, and other slide navigation functions.

```
graph LR; A[Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages] --> B[Background]; A --> C[Runtime Scheduling]; B --> D[Background: Runtime Scheduling];
```

2014-08-11

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

- Background
- Runtime Scheduling
- Background: Runtime Scheduling

Background: Runtime Scheduling

- Schedulers can be defined in a discrete manner:
 - Queue a process into task
 - Select a task
 - Execute process until it dies
- Schedulers can be gathered in every step into queues
 - Number of input entries
 - Ordering of task sets
 - Number of children sets
- Other statistics are necessary for recognizing cooperativity?

```
graph LR; A[Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages] --> B[Background]; A --> C[Runtime Scheduling]; B --> D[Background: Runtime Scheduling];
```

2014-08-11

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

- Background
- Runtime Scheduling
- Background: Runtime Scheduling

Background: Runtime Scheduling

- Schedulers can be defined in a discrete manner:
 - Queue a process into task
 - Select a task
 - Execute process until it dies
- Schedulers can be gathered in every step into queues
 - Number of input entries
 - Ordering of task sets
 - Number of children sets
- Other statistics are necessary for recognizing cooperativity?

- 2014-08-11

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

 - └ Background
 - └ Runtime Scheduling
 - └ Background: Runtime Scheduling

Background: Runtime Scheduling

 - Scheduler can be defined in a discrete manner
 - Choose a partitioning rule
 - Choose a scheduling policy
 - Update queue scheduling state
 - Schedulers can be gathered at every step about progress
 - Number of shared partners
 - Number of local ops
 - Number of messages, etc.
 - Other policies are necessary for recognizing cooperativity?

- Choosing from set: Top Always (batching), Queue (Round-Robin)
 - Reductions Return some indication: yield/blocked, unblocked, completed, nothing
 - Update state based on historical data.
 - Timestamp of last run = Separate between batching/round-robin
 - Cooperativity metrics: yields, partners, thus longevity and granularity
 - gets back to why we chose swap:
 - (if asymmetrical, we would have a harder time with 'partner')
 - (if async, we would have a harder time with 'yield' and noticing longevity/progress)

Current Section:

2014-08-11	Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages	Current Section
	└ ErLam	
	└ Current Section:	
		■ Erlang
		■ The Language
		■ Channel Implementations
		■ Processes & Virtualization

2014-08-11	Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages	Current Section
	└ ErLam	
	└ Current Section:	
		■ Erlang
		■ The Language
		■ Channel Implementations
		■ Processes & Virtualization

- Current Section:

2 Erlam

 - The Language
 - Channel Implementations
 - Simulation & Visualization

Navigation icons: back, forward, search, etc.

- Current Section:

2014-08-11

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

ErLam

Current Section:

Current Section

■ ErLam

■ The Language
■ Channel Implementations
■ Simulation & Visualization

2 ErLam

■ The Language

■ Channel Implementations

■ Simulation & Visualization

◀ ◻ ▶ ◀ ☰ ▶ ◀ ⌂ ▶ ◀ 🔍 ▶

ErLam: The Language

[illegible]

```
graph TD; A[Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages] --> B[ErLam]; A --> C[The Language]; B --> D[ErLam: The Language]
```

The diagram illustrates the relationship between different components of the research. At the top level is "Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages". This branches down to "ErLam" and "The Language". "ErLam" further branches down to "ErLam: The Language".

[illegible]

ErLam: The Language

```
<Expression> ::= <Variable>
                | <Integer>
                | 'newchan'
                | '(' <Expression> ')',
                | <Expression> <Expression>
                | 'if' <Expression> <Expression> <Expression>
                | 'swap' <Expression> <Expression>
                | 'spawn' <Expression>
                | 'fun' <Variable> '.' <Expression>
```



- 2014-08-11

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

ErLam

The Language

ErLam: The Language

- Extremely simple on purpose (5 keywords).
 - Issue now began to be how to build up test primitives
 - Made a library which allowed for built ins.

```
<expression> ::= <operation>
               | <constant>
               | <variable>
               | <lambda>
               | <let> <expression> <expression>
               | <let*> <expression> <expression>
               | <letrec> <expression> <expression>
               | <letrec*> <expression> <expression>
               | <letrec*letrec*> <expression> <expression>
               | <letrec*letrec*letrec*> <expression> <expression>
               | <letrec*letrec*letrec*letrec*> <expression> <expression>
```

ErLam: The Language

```
elib
// ...
omega = (fun x.(x x));
// ...
add = _erl[2]{ fun(X) when is_integer(X) ->
                fun(Y) when is_integer(Y) ->
                    X+Y
                end
            end
        };
// ...
bile
```

Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-11

ErLam
The Language
ErLam: The Language

ErLam: The Language

```

// ...
omega = (fun x.(x x));
// ...
add = _erl[2]{ fun(X) when is_integer(X) ->
                fun(Y) when is_integer(Y) ->
                    X+Y
                end
            end
        };
// ...
bile
```

- There's options for built-ins as well as macros.
- Built-ins are raw Erlang, gets wrapped up into AST, and still "reduces" the same (i.e. no multi-variable functions).

ErLam: The Language

Example Application: Parallel Fibonacci

```
// pfib.els -
fun N.
    (omega fun f,m.(
        if (leq m 1)
            m
            (merge fun _.(f f (sub m 1))
                    fun _.(f f (sub m 2))
                    add)) N)
```

```
$ els pfib.els      (Compile the script)
$ ./pfib.ex -r 10  (Finds the 10th Fibonacci number)
```

Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-11

ErLam
The Language
ErLam: The Language


ErLam: The Language

```

// pfib.els -
fun N.
    (omega fun f,m.(
        if (leq m 1)
            m
            (merge fun _.(f f (sub m 1))
                    fun _.(f f (sub m 2))
                    add)) N)
```

- R option is to run program applied to 10.
- To add more to this presentation than just reading paper I want to also give more detail as to system usage.
- Explain this Common Usage Pattern.

ErLam: Channel Implementations




Process Blocking Swap

Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages


2014-08-11

ErLam
Channel Implementations
ErLam: Channel Implementations



- Blocking: Maintains state of current and previous swap value until swap is completed.
- Mention expected effects on scheduler.
 - Scheduler will keep hold of the process, needs to recheck if blocked.
 - If all communicating, large process queue of blocked processes.
- Blocking is default, passing '-a' to els for absorb

ErLam: Channel Implementations




Process Absorption Swap

Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-11

ErLam
Channel Implementations
ErLam: Channel Implementations



Absorption swap is visually easier to recognize.

1. Whole process gets absorbed by channel, away from scheduler.
 2. When the second process completes the swap, the process gets removed from channel.
 - The $p0$ process can go back to its original scheduler,
 - OR to scheduler which unblocked it.
- Effects on scheduler are obvious.

ErLam: Simulation & Visualization

- System Behaviours:
 - Degree of Parallelism
 - Consistency of Cooperation
 - Degree of Longevity/Interactivity
 - Partial System Cooperativity
- Logging & Report Generation

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-11

ErLam
Simulation & Visualization
ErLam: Simulation & Visualization



- We first want to look at four types of system behaviour ranges:
 - Parallelism: Gets back to Ring vs Cloud, Compare the two.
 - Consistency: Ring/Star=consistent, but if given a random choice.
 - Longevity: Ratio of Communicating/Computing processes.
 - Longevity = Time spent reducing (low=communicator)
 - Interactivity, also takes into account user interaction.
 - Full vs Partial Cooperation: Multiple groups of Stars or Rings.
- Definitely won't get to all behaviour tests that were in report, but will focus on the key ones for each scheduler mechanic.
- Finally, quickly, go over the current report generation that ErLam can perform.

Navigation icons: back, forward, search, etc.

ErLam: Simulation & Visualization

System Behaviours: Degree of Parallelism





Figure: $PRing_N$ and $ClusterComm_{(N,1)}$ primitives to test degree of parallelism.

Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-11

ErLam
Simulation & Visualization
ErLam: Simulation & Visualization



- Brings back the Ring and Star.
- We call the left, $PRing$, with the parameter N = number of processes.
- We call the right, $ClusterComm$ with two parameters, N like $PRing$, $M = 1$ in this cas

ErLam: Simulation & Visualization

System Behaviours: Consistency of Cooperation





Figure: $ClusterComm_{(N,M)}$ to test effect of consistency on scheduler.

Navigation icons: back, forward, search, etc.

2014-08-11
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
ErLam
Simulation & Visualization
ErLam: Simulation & Visualization



- We can vary number of channels in relation to processes to check the effects of inconsistency on Cooperative-Conscious schedulers.
- Worst case scenario for C-C schedulers.

ErLam: Simulation & Visualization

System Behaviours: Degree of Interactivity





Figure: $UserInput_{(T,C)}$, simulates user interaction or a number (C) of external/timed (T) events.

Navigation icons: back, forward, search, etc.

2014-08-11
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
ErLam
Simulation & Visualization
ErLam: Simulation & Visualization



- To test longevity we can just vary the length of time the processes chug for all tests.
- To test interactivity though, we need a way to simulate user interaction.
- $UserInput$ captures hanging for a single event. We can compose these: $< NEXT >$
- With our cloud of processes (also called chugmachine) for simulating a program with consistent working processes and processes which are interactive.

ErLam: Simulation & Visualization

System Behaviours: Degree of Interactivity





Figure: $Interactivity_{(N,M)}$, composure of $ChugMachine_N$ (Cloud), and M instances of $UserInput_{(5,2)}$.

Navigation icons: back, forward, search, etc.

2014-08-11
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
ErLam
Simulation & Visualization
ErLam: Simulation & Visualization



- To test longevity we can just vary the length of time the processes chug for all tests.
- To test interactivity though, we need a way to simulate user interaction.
- $UserInput$ captures hanging for a single event. We can compose these: $< NEXT >$
- With our cloud of processes (also called chugmachine) for simulating a program with consistent working processes and processes which are interactive.

ErLam: Simulation & Visualization

System Behaviours: Partial System Cooperativity





Figure: $PTree_{(W,N)}$, a composure of W $ClusterComm_{(N,1)}$ instances running concurrently.

Navigation icons: back, forward, search, etc.

2014-08-11
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
ErLam
Simulation & Visualization
ErLam: Simulation & Visualization



- Previously (besides Interactivity), all systems were full system cooperation.
- We can of course use Interactivity for our partial system cooperativity tests, but we would instead like to see logical grouping.
- Hence the set of Work groups (or stars).

ErLam: Simulation & Visualization

Logging & Report Generation

Things we could log:

Navigation icons: back, forward, search, etc.

2014-08-11
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
ErLam
Simulation & Visualization
ErLam: Simulation & Visualization



ErLam: Simulation & Visualization

Logging & Report Generation

Things we could log:

- Process Queue Size (per LPU)

Navigation icons: back, forward, search, etc.

2014-08-11
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
ErLam
Simulation & Visualization
ErLam: Simulation & Visualization



- Queue-Length: work-stealing mechanics and saturation ability.

ErLam: Simulation & Visualization


Things we could log:

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts



Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-11



Things we could log:

- Business Queue Size

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts

- Queue-Length: work-stealing mechanics and saturation ability.
- Tick-Action: Visualize the density of computation/communication.

ErLam: Simulation & Visualization


Things we could log:

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts
- State of the Scheduler (waiting/running)



Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-11



Things we could log:

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Presamples
- State of the Scheduler (waiting/running)

- Queue-Length: work-stealing mechanics and saturation ability.
- Tick-Action: Visualize the density of computation/communication.
- Sched-State: Useful for comparing stealing/process selection mechanics.

ErLam: Simulation & Visualization


Things we could log:

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts
- State of the Scheduler (waiting/running)
- Channel State (Blocked/Unblocked)



Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-11



Things we could log:

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts
- State of the Scheduler (waiting/running)

- Queue-Length: work-stealing mechanics and saturation ability.
- Tick-Action: Visualize the density of computation/communication.
- Sched-State: Useful for comparing stealing/process selection mechanics.
- Chan-State: Tracking interactivity, speed of unblock=attentive to cooperation.

ErLam: Simulation & Visualization

Things we could log:

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts
- State of the Scheduler (waiting/running)
- Channel State (Blocked/Unblocked)
- ...

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-11

- ErLam
 - └Simulation & Visualization
 - └ErLam: Simulation & Visualization

ErLam: Simulation & Visualization

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts
- State of the Scheduler (waiting/running)
- Channel State (Blocked/Unblocked)
-

- Queue-Length: work-stealing mechanics and saturation ability.
- Tick-Action: Visualize the density of computation/communication.
- Sched-State: Useful for comparing stealing/process selection mechanics.
- Chan-State: Tracking interactivity, speed of unblock=attentive to cooperation.
- Of course there are more, but we limited ourselves to the above for initial testing purposes.

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻

ErLam: Simulation & Visualization

Things we could log:

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts
- State of the Scheduler (waiting/running)
- Channel State (Blocked/Unblocked)
- ...

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-11

- ErLam
 - └Simulation & Visualization
 - └ErLam: Simulation & Visualization

- Things we could log:
 - Process Queue Size (per LPU)
 - Quantity of Reductions/Yields/Presumptive
 - State of the Scheduler (waiting/running)
 - Channel State (Blocked/Unblocked)

- Queue-Length: work-stealing mechanics and saturation ability.
- Tick-Action: Visualize the density of computation/communication.
- Sched-State: Useful for comparing stealing/process selection mechanics.
- Chan-State: Tracking interactivity, speed of unblock=attentive to cooperation.
- Of course there are more, but we limited ourselves to the above for initial testing purposes.

◀ ◻ ▶ ◀ ☰ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺

Current Section:

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-11

- └ Scheduler Implementations
 - └ Current Section:

Current Section

- Feedback-based Scheduling
- Example Schedulers
- Feedback Scheduler

- 3 Scheduler Implementations

◀ ◻ ▶ ◀ ▢ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ↺ 🔍 ↻

Scheduler Implementations: Example Schedulers

- Single-Thread Dual-Queue
 - Translation of CML scheduler
- Round-Robin with Single Global Queue (MTRRGQ)
 - No Work-Stealing
 - # LPUs can vary
- Round-Robin with Work-Stealing (MTRRWS)
 - Steal via direct access
 - Steal via advertisement

- First tested translating a known scheduler to the discrete step scheduler semantics.
- MTRRGQ: all synchronization around a single global queue, test $P=1$ or max.
- MTRRWS-SQ:
 - Schedulers can access a random LPU's queue end and steal from their bottoms.
- MTRRWS-IS:
 - Uses a secondary queue to advertise desire to steal.
 - Scheduler can check secondary queue as desired.
 - No requirement for synchronization on private process queue.

Scheduler Implementations: Feedback Schedulers

Three types of mechanics:

- Longevity-Based Batching
- Channel Pinning
- Bipartite-Graph Aided Sorting

- Instead of a single cooperativity-conscious scheduler, we implemented three mechanics which take cooperativity into account on top of the basic schedulers.

Scheduler Implementations: Feedback Schedulers

Longevity-Based Batching

- Choose via Round-Robin
 - from batch rather than queue
 - keeps track of number of rounds (batch size)
- Work-Steal whole batches
- Spawn to batch unless: $|b_i| \geq B$
 - Make singleton with new process.
 - Push parent and child into new batch.

- Batching processes based on longevity.
 - Based on $\text{occam-}\Pi$.
 - if a process communicates frequently then it will be batched (absorption), singleton if very computation-bound.
- We are normal RR but with one extra layer.
- If batch is too big during spawns we can:
 - Make singleton, best if child is needed to start work right away. Map-Reduce.
 - Make push-back, parent can get another chance to spawn more children sooner.

Scheduler Implementations: Feedback Schedulers

Longevity-Based Batching

- Choose via Round-Robin
 - from batch rather than queue
 - keeps track of number of rounds (batch size)
- Work-Steal whole batches
- Spawn to batch unless: $|b_i| \geq B$
 - Make singleton with new process.
 - Push parent and child into new batch.

GOAL: Can batching based on longevity account for fine/coarse parallelism in application?

Navigation icons: back, forward, search, etc.

2014-08-11
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Scheduler Implementations
Feedback Schedulers
Scheduler Implementations: Feedback Schedulers

Scheduler Implementations: Feedback Schedulers
Longevity-Based Batching

- Choose via Round-Robin
 - from batch rather than queue
 - keeps track of number of rounds (batch size)
- Work-Steal whole batches
- Spawn to batch unless: $|b_i| \geq B$
 - Make singleton with new process.
 - Push parent and child into new batch.

GOAL: Can batching based on longevity account for fine/coarse parallelism in application?

- Batching processes based on longevity.
 - Based on occam-Π.
 - if a process communicates frequently then it will be batched (absorption), singleton if very computation-bound.
- We are normal RR but with one extra layer.
- If batch is too big during spawns we can:
 - Make singleton, best if child is needed to start work right away. Map-Reduce.
 - Make push-back, parent can get another chance to spawn more children sooner.

Scheduler Implementations: Feedback Schedulers

Channel-Pinning

- Upon call to *newchan*, pin to LPU based on spread algorithm:
 - *same* - LPU *newchan* is called is where it is pinned.
 - *even* - Cycle through LPUs and pin based on that.
 - ...
- Work-steal based on channel that's been pinned to you.

Navigation icons: back, forward, search, etc.

2014-08-11
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Scheduler Implementations
Feedback Schedulers
Scheduler Implementations: Feedback Schedulers

Scheduler Implementations: Feedback Schedulers
Channel-Pinning

- Upon call to *newchan*, pin to LPU based on spread algorithm
 - *same* - LPU *newchan* is called is where it is pinned.
 - *even* - Cycle through LPUs and pin based on that.
- Work-steal based on channel that's been pinned to you.

GOAL: Can an even-like spread increase early saturation?

- Pin channels to LPUs.
 - Pinning a channel means to set a process affinity to a LPU based on the channels it uses.
 - Work-Stealing works like Go-Fish.

Scheduler Implementations: Feedback Schedulers

Channel-Pinning

- Upon call to *newchan*, pin to LPU based on spread algorithm:
 - *same* - LPU *newchan* is called is where it is pinned.
 - *even* - Cycle through LPUs and pin based on that.
 - ...
- Work-steal based on channel that's been pinned to you.

GOAL: Can an *even*-like spread increase early saturation?

Navigation icons: back, forward, search, etc.

2014-08-11
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Scheduler Implementations
Feedback Schedulers
Scheduler Implementations: Feedback Schedulers

Scheduler Implementations: Feedback Schedulers
Channel-Pinning

- Upon call to *newchan*, pin to LPU based on spread algorithm
 - *same* - LPU *newchan* is called is where it is pinned.
 - *even* - Cycle through LPUs and pin based on that.
- Work-steal based on channel that's been pinned to you.


GOAL: Can an even-like spread increase early saturation?

- Pin channels to LPUs.
 - Pinning a channel means to set a process affinity to a LPU based on the channels it uses.
 - Work-Stealing works like Go-Fish.

Scheduler Implementations: Feedback Schedulers

Bipartite-Graph Aided Sorting

- Based on Round-Robin & Work-stealing
- Keep track of events which may effect cooperativity:
 - Spawning
 - Blocking/Unblocking
 - Steals
- If number of events over some threshold, re-sort.



2014-08-11

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

- └ Scheduler Implementations
 - └ Feedback Schedulers
 - └ Scheduler Implementations: Feedback Schedulers

Scheduler Implementations: Feedback Schedulers

2014-08-11


- Based on Round-Robin & Work-stealing
- Keep track of events which may effect cooperativity
 - Spawning
 - Blocking/Unblocking
 - Steals
- If number of events over some threshold, re-sort.

- Keep a list of all communications as a graph between set of processes and channels.

Scheduler Implementations: Feedback Schedulers

Bipartite-Graph Aided Sorting

- Based on Round-Robin & Work-stealing
- Keep track of events which may effect cooperativity:
 - Spawning
 - Blocking/Unblocking
 - Steals
- If number of events over some threshold, re-sort.



2014-08-11

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

- └ Scheduler Implementations
 - └ Feedback Schedulers
 - └ Scheduler Implementations: Feedback Schedulers

Scheduler Implementations: Feedback Schedulers

2014-08-11

- Based on Round-Robin & Work-stealing
- Keep track of events which may effect cooperativity
 - Spawning
 - Blocking/Unblocking
 - Steals
- If number of events over some threshold, re-sort.

- Keep a list of all communications as a graph between set of processes and channels.

GOAL: Are alternate channel implementations worth exploration?

Current Section:

2014-08-11

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

- └ Results
 - └ Current Section:

Current Section

Results

Results:

Longevity-Based Batching

- Can batching based on longevity recognize fine/coarse parallelism in an application?

Channel Pinning

- Can an even-like spread increase early saturation?

Bipartite-Graph Aided Sorting

- Are alternate channel implementations worth exploration?

Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-11

Results

Results:

- Remind about the goals of the talk. However there are other things we can of course study now that we have a simulator:
- LBB: Can also look at how the batch size effects different types of behaviours.
- CP: Could also look at a stealing mechanism.

Results

- Longevity-Based Batching
 - Can batching based on longevity recognize fine/coarse parallelism in an application?
- Channel Pinning
 - Can an even-like spread increase early saturation?
- Bipartite-Graph Aided Sorting
 - Are alternate channel implementations worth exploration?

Results: Longevity-Based Batching




Table: Comparison of different sized $PRing_N$ on the Longevity Batching Scheduler with batch size $B = 10$.

Navigation icons: back, forward, search, etc.


Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-11

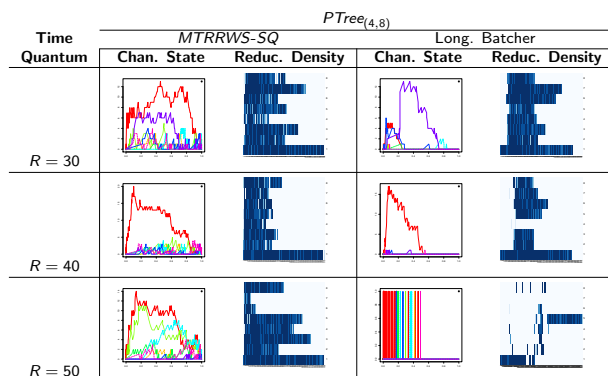
Results

Results: Longevity-Based Batching

- Early tests gave promising results.
- Here is $PRing_N$ which shows the reabsorption and containment on a single LPU as expected and hoped.
- So does batching based on longevity really recognize fine/coarse parallelism in an application?
- Sort of, if you know what the right time-quantum is to make that distinction.



Results: Longevity-Based Batching



Navigation icons: back, forward, search, etc.


Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-11

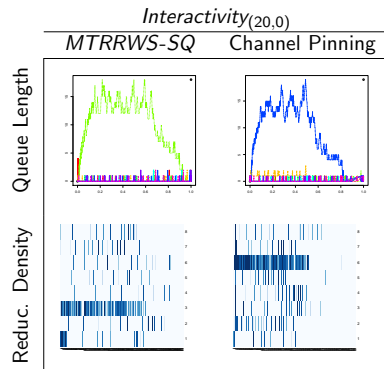
Results

Results: Longevity-Based Batching

- Comparison of $PTree_{(4,8)}$ running with the Longevity-Based Batching Scheduler and $MTRRWS-SQ$ at different time-quantums.
- At lower time quantum Long. Batcher starts to look like $RRWS-SQ$, however batching and absorption channels tend to lead to consolidation.
- At higher time quantum Long. Batcher results in the originally expected work-groups. But it turns out to be inefficient due to lost chances of parallelism of each "star" of each group.
- Heuristical adjustment of the time-quantum would definitely be possible.



Results: Channel Pinning




Navigation icons: back, forward, search, etc.

2014-08-11 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

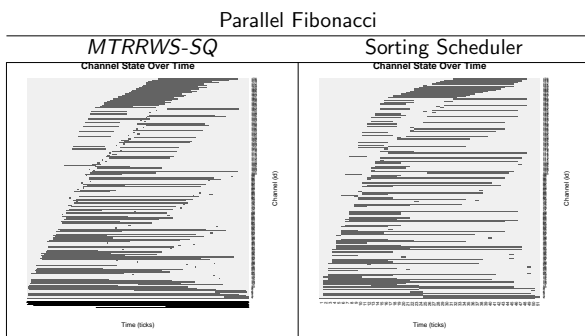
Results

Results: Channel Pinning



- Comparison of Uniform synchronization for *MTRRWS-SQ* and the Channel Pinning Scheduler on Absorption Channels.
- This used the *even* spread type.
- Note the speed at which it saturates all cores.
- Despite Naive WS, we still have decent spread.

Results: Bipartite-Aided Graph Sorting




Navigation icons: back, forward, search, etc.

2014-08-11 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Results

Results: Bipartite-Aided Graph Sorting



- Had to deviate from the primitives. No primitive relied on process order.
- PFib has a strong reliance on order of execution.
- Channel State comparison of Parallel Fibonacci executed on *MTRRWS-SQ* and the Bipartite-Graph Aided Sorting Scheduler.
- Note the large reduction in number of ticks.

Current Section:

2014-08-11 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Conclusions & Future Work

Current Section:



5 Conclusions & Future Work

- ErLam Toolkit
- Cooperative Schedulers
- Cooperativity as a Metric

Navigation icons: back, forward, search, etc.

Conclusions & Future Work: ErLam Toolkit

- ◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ 🔍 ↺



-

- ## Conclusions & Future Work: Cooperative Scheduling
- **Longevity Batching:**
 - Visual benefit from heuristic based Quantum selection.
 - At 2 stands, limited gain from longevity recognition.
 - **Channel Pacing:**
 - Promising saturation and queue-stretching analysis.
 - **Slipstream-Graph Aided Scheduling:**
 - Supporting results on MapReduce style applications.
 - Worth studying blocking Channels further for gains from

- Possible to recognize and benefit from.
- The three example mechanisms are promising and can be extended for practical modern languages, despite simplistic simulation language.
- More to explore:
 - Alternate Message-Passing Types (Asymmetric?)

-

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻

< □ > < ▢ > < ≡ > < ≡ > ≡ ↺ 🔍 ↻