

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Alexander Dean

Rochester Institute of Technology
Golisano College of Computer and Information Sciences

August 12, 2014

Navigation icons

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
Alexander Dean
Rochester Institute of Technology
Golisano College of Computer and Information Sciences
August 12, 2014

- Thank Fluet, Heliotis, and Raj.
- Dedicate to parents, who are unable to be present.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

└ Process Cooperativity as a Feedback Metric

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
1 Background
2 Runtime Scheduling
3 Cooperativity
4 Message Passing
5 Scheduler Implementations
6 Example Schedulers
7 Feedback Mechanisms
8 Results
9 Conclusions & Future Work
10 ErLam Toolkit
11 The Language
12 Channel Implementations
13 Simulation & Visualization
14 Cooperativity as a Metric

- 1 Background
 - Runtime Scheduling
 - Cooperativity
 - Message Passing
- 2 ErLam Toolkit
 - The Language
 - Channel Implementations
 - Simulation & Visualization
- 3 Scheduler Implementations
 - Example Schedulers
 - Feedback Mechanisms
- 4 Results
- 5 Conclusions & Future Work
 - ErLam Toolkit
 - Cooperative Schedulers
 - Cooperativity as a Metric

Navigation icons

- Mouthful of a title, so I'll break it up:
 1. Runtime scheduling, to give some grounding in the area of study.
 2. Cooperativity, what it is and motivation to use it.
 3. Message Passing, because, as it turns out, it's a nice abstraction for our purpose of capturing cooperativity.
- The core of the work revolves around the toolkit I built.
 - A language/compiler/runtime/testing-framework
 - But also a *Simulator* which has a plug-and-play scheduler API. It let me test schedulers on a common test bed.
- Next, go over the list of schedulers & feedback mechanisms.
- Results, Conclusions, & Future Work.

Background

2014-08-12

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

└ Background

└ Background

Background
1 Background
2 Runtime Scheduling
3 Cooperativity
4 Message Passing

Introduce the new section:

1 Background

- Runtime Scheduling
- Cooperativity
- Message Passing

- 1 Background
 - Runtime Scheduling
 - Cooperativity
 - Message Passing

Navigation icons

Background: Runtime Scheduling

- Schedulers can be defined in a discrete manner:
 - 1 Choose a process from set,
 - 2 Reduce it,
 - 3 Update private scheduler state.

Background: Runtime Scheduling

- Schedulers can be defined in a discrete manner:
 - 1 Choose a process from set,
 - 2 Reduce it,
 - 3 Update private scheduler state.
- Statistics can be gathered at every step about process:
 - Timestamp of last run,
 - Number of reductions, *etc.*

Background: Runtime Scheduling

- Schedulers can be defined in a discrete manner:
 - 1 Choose a process from set,
 - 2 Reduce it,
 - 3 Update private scheduler state.
- Statistics can be gathered at every step about process:
 - Timestamp of last run,
 - Number of reductions, etc.
- *What statistics are useful?*

2014-08-12

Background: Runtime Scheduling

- We can look at process schedulers like a function:
 - Takes a set of processes, and some private state.
 - Job of the function is to choose a process, and run it for a bit.
 - Then, based on what happened while running process, we update the state.
- Big questions: How are we choosing a process? What should effect our decision?

2014-08-12

Background: Runtime Scheduling

- We can look at process schedulers like a function:
 - Takes a set of processes, and some private state.
 - Job of the function is to choose a process, and run it for a bit.
 - Then, based on what happened while running process, we update the state.
- Big questions: How are we choosing a process? What should effect our decision?
- Timestamp of last run? →
 - Choose always most recent, it's a batch scheduler.
 - Choose oldest, we get something called Round-Robin.
- Number of reductions? → longevity = might want to give someone else a go.
- What are useful, and what do they tell us about the state of the system? Well this leads us to cooperativity.

2014-08-12

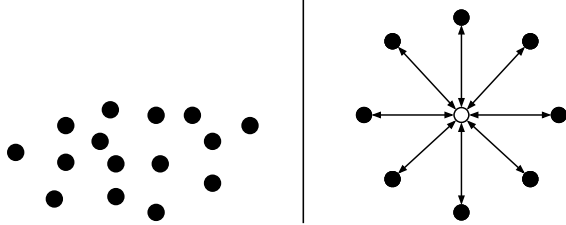
References

- Schedulers can be defined in a discrete manner:
 - Choose a process from set,
 - Execute it,
 - Update private scheduler state.
- Statistics can be gathered at every step about process:
 - Timestamp of last run,
 - Number of reductions, etc.
- Other statistics are useful?

- We can look at process schedulers like a function:
 - Takes a set of processes, and some private state.
 - Job of the function is to choose a process, and run it for a bit.
 - Then, based on what happened while running process, we update the state.
- Big questions: How are we choosing a process? What should effect our decision?
- Timestamp of last run? →
 - Choose always most recent, it's a batch scheduler.
 - Choose oldest, we get something called Round-Robin.
- Number of reductions? → longevity = might want to give someone else a go.
- What are useful, and what do they tell us about the state of the system? Well this leads us to cooperativity.

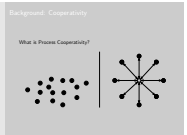
Background: Cooperativity

What is Process Cooperativity?



Navigation icons: back, forward, search, etc.

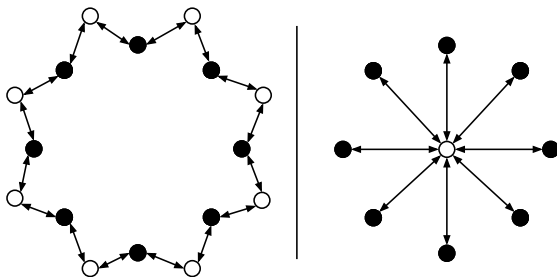
2014-08-12
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
└ Background
└ Cooperativity
└ Background: Cooperativity



- What is Process Cooperativity?
- White = channel & black = a process.
- Can think of channel a mechanism for passing information between processes.
 - These are nice functional abstractions of things like locks, shared-memory, etc.
- Left: Cloud of processes with no interaction.
- Right: We see a definite structure caused by some sharing of information. This is the core of recognizing cooperation, namely, recognizing these structures when they exist.

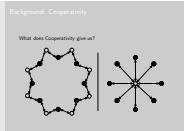
Background: Cooperativity

What does Cooperativity give us?



Navigation icons: back, forward, search, etc.

2014-08-12
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
└ Background
└ Cooperativity
└ Background: Cooperativity



What does Cooperativity give us? What's the difference in the behaviour of cooperation in the left/right applications?

- Left: A Ring,
 - the level of parallelism is nearly nil.
 - Each process is cooperating yes, but granularity is very fine.
- Right: A Star,
 - the level of parallelism is nearly full.
 - Each process is cooperating, **not reliant on more than one** other process.
- In both, the whole system is communicating, but with cooperation, we can find the level of parallelism possible.

Next: Knowing this, how can we recognize cooperativity? Seems to be all about recording interactions with the channel.

Background: Message Passing

We use a Symmetric, Synchronous, Message-Passing Primitive:

swap

- Purely captures cooperation of processes through synchronizing on a shared channel.

Navigation icons: back, forward, search, etc.

2014-08-12
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
└ Background
└ Message Passing
└ Background: Message Passing



- Symmetric, Synchronous, Message-Passing primitive.
- Symmetric:
 - Only one message passing primitive: SWAP
- Synchronous:
 - Blocks until it's partner gets there.
- Purely captures cooperation: Simple synchronization representation.
- This is really what I based the language on.
- So, what does the rest of the language look like.

ErLam Toolkit

2 ErLam Toolkit

- The Language
- Channel Implementations
- Simulation & Visualization

2014-08-12
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

└ ErLam Toolkit

└┐ ErLam Toolkit

ErLam Toolkit

- The Language
- Channel Implementations
- Simulation & Visualization

Introduce the new section:

- 2 ErLam Toolkit
 - The Language
 - Channel Implementations
 - Simulation & Visualization

ErLam Toolkit: The Language

```
<Expression> ::= <Variable>
| <Integer>
| 'newchan'
| '(' <Expression> ')'
| <Expression> <Expression>
| 'if' <Expression> <Expression> <Expression>
| 'swap' <Expression> <Expression>
| 'spawn' <Expression>
| 'fun' <Variable> '.' <Expression>
```

2014-08-12
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

└ ErLam Toolkit

└ The Language

└┐ ErLam Toolkit: The Language

ErLam Toolkit: The Language

```
is_integer :: Integer -> Boolean
is_integer = \x -> (x < 0 || x > 0)

newchan :: IO Channel
newchan = do { h <- newHandle 1024; return h }

if :: Expression -> Expression -> Expression -> Expression
if e1 e2 e3 = do { h <- newHandle 1024; return h }

spawn :: Expression -> Expression
spawn e = do { h <- newHandle 1024; return h }

fun :: Variable -> Expression -> Expression
fun v e = do { h <- newHandle 1024; return h }
```

- Extremely simple on purpose (5 keywords).
- Issue now began to be how to build up primitive test behaviours.
- Made a library which allowed for built ins.

ErLam Toolkit: The Language

```
elib
// ...
ignore = (fun _.(fun y.y));
omega = (fun x.(x x));
// ...
add = _erl[2]{ fun(X) when is_integer(X) ->
                fun(Y) when is_integer(Y) ->
                    X+Y
                end
            end
};

// ...
bile
```

2014-08-12
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

└ ErLam Toolkit

└ The Language

└┐ ErLam Toolkit: The Language

ErLam Toolkit: The Language

```
is_integer :: Integer -> Boolean
is_integer = \x -> (x < 0 || x > 0)

newchan :: IO Channel
newchan = do { h <- newHandle 1024; return h }

if :: Expression -> Expression -> Expression -> Expression
if e1 e2 e3 = do { h <- newHandle 1024; return h }

spawn :: Expression -> Expression
spawn e = do { h <- newHandle 1024; return h }

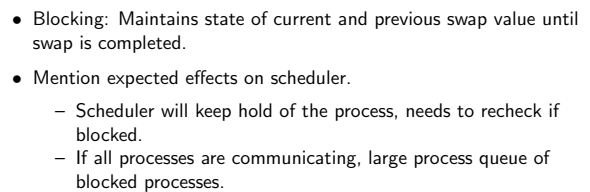
fun :: Variable -> Expression -> Expression
fun v e = do { h <- newHandle 1024; return h }
```

- There's options for built-ins as well as macros.
- Built-ins are raw Erlang, gets wrapped up into AST, and still "reduces" the same (i.e. no multi-variable functions).

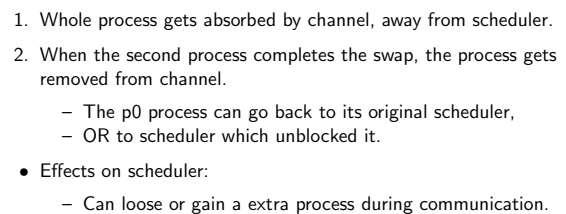
Example Application: Simple Swap

- Here is a simple application which:
 - Spawns a process to swap the number 42
 - Calls swap to get the value from the other process.
- Build up from here to simulate more complex behaviour.
- We can "do some work" before swaping, *etc.*
- This is how we built up our primitive test behaviours.

Process Blocking Swap



Process Absorption Swap



ErLam Toolkit: Simulation & Visualization

Primitive Testing Behaviours:

- Degree of Parallelism
- Partial System Cooperativity

Logging & Report Generation

2014-08-12
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
ErLam Toolkit
Simulation & Visualization
ErLam Toolkit: Simulation & Visualization

ErLam Toolkit: Simulation & Visualization
Primitive Testing Behaviours:
■ Degree of Parallelism
■ Partial System Cooperativity
Logging & Report Generation

- To test cooperativity, we need a set of test behaviours.
- We want to analyze a large number of structure types.
- Two of them that I studied are:
 - Parallelism: Gets back to Ring vs Star, compare the two.
 - Full vs Partial Cooperation: Multiple groups of Stars or Rings.
- Finally, quickly, go over the current report generation that ErLam performs.

Navigation icons: back, forward, search, etc.

ErLam Toolkit: Simulation & Visualization

System Behaviours: Degree of Parallelism




Figure: $PRing_N$, and $ClusterComm_{(N,1)}$ primitives to test degree of parallelism.

Navigation icons: back, forward, search, etc.

2014-08-12
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
ErLam Toolkit
Simulation & Visualization
ErLam Toolkit: Simulation & Visualization

ErLam Toolkit: Simulation & Visualization
Figure: $PRing_N$ and $ClusterComm_{N,1}$ primitives to test degree of parallelism.

- Here are our Ring and Star behaviours, we call them by different names.
- We call the left, $PRing$, with the parameter N = number of processes.
- We call the right, $ClusterComm$ with two parameters, N like $PRing$, $M = 1$ in this case.
- Note the generalization of the Star is a set of M channels that are randomly accessed.

ErLam Toolkit: Simulation & Visualization

System Behaviours: Consistency of Cooperation

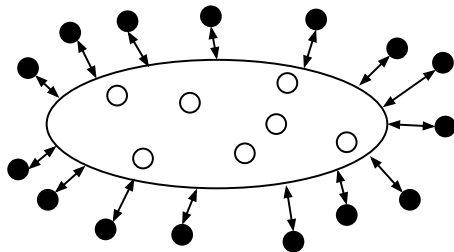


Figure: $ClusterComm_{(N,M)}$ to test effect of consistency on scheduler.

Navigation icons: back, forward, search, etc.

2014-08-12
Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
ErLam Toolkit
Simulation & Visualization
ErLam Toolkit: Simulation & Visualization

ErLam Toolkit: Simulation & Visualization
Figure: $ClusterComm_{N,M}$ to test effect of consistency on scheduler.

- Worst case scenario for C-C schedulers.
- However, this is still full system cooperativity. We want some minor work groups.

ErLam Toolkit: Simulation & Visualization

System Behaviours: Partial System Cooperativity

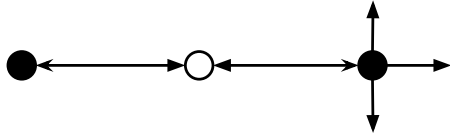


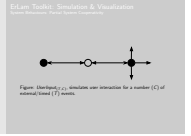
Figure: $UserInput_{(T,C)}$, simulates user interaction for a number (C) of external/timed (T) events.

Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12

ErLam Toolkit
Simulation & Visualization
ErLam Toolkit: Simulation & Visualization



- So we can make tiny work groups of two processes, like our simple swapping program earlier.
- The difference here is minor, we call this test UserInput because one process can hang for a set amount of time before triggering an event (effectively simulating user interaction).
- But we need multiple of these work groups, running concurrently.
- Run a set of them in parallel, and we could even inject a bit of overhead at will by running any number of extra processes.

ErLam Toolkit: Simulation & Visualization

System Behaviours: Partial System Cooperativity

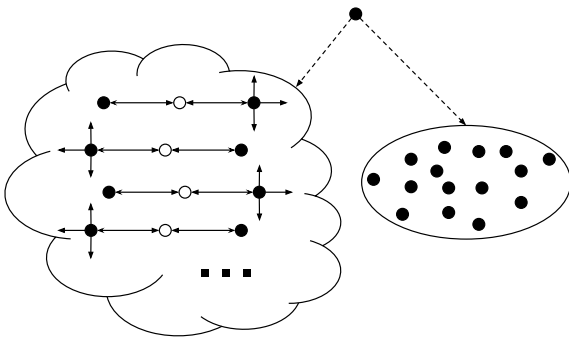


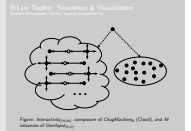
Figure: $Interactivity_{(N,M)}$, composure of $ChugMachine_N$ (Cloud), and M instances of $UserInput_{(5,2)}$.

Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12

ErLam Toolkit
Simulation & Visualization
ErLam Toolkit: Simulation & Visualization



- So we can make tiny work groups of two processes, like our simple swapping program earlier.
- The difference here is minor, we call this test UserInput because one process can hang for a set amount of time before triggering an event (effectively simulating user interaction).
- But we need multiple of these work groups, running concurrently.
- Run a set of them in parallel, and we could even inject a bit of overhead at will by running any number of extra processes.

ErLam Toolkit: Simulation & Visualization

System Behaviours: Partial System Cooperativity

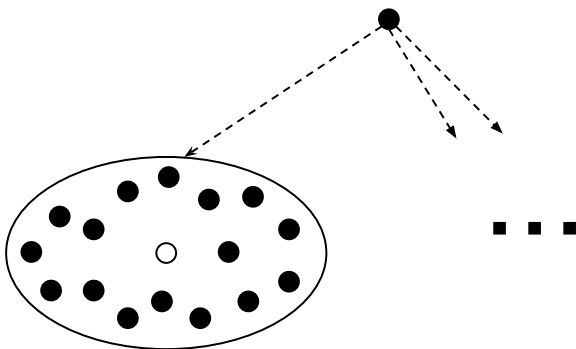


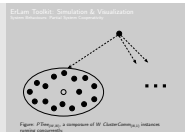
Figure: $PTree_{(W,N)}$, a composure of W $ClusterComm_{(N,1)}$ instances running concurrently.

Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12

ErLam Toolkit
Simulation & Visualization
ErLam Toolkit: Simulation & Visualization



- But we can run ClusterComm or PRing in parallel.

ErLam Toolkit: Simulation & Visualization

Logging & Report Generation

Things we could log:

- Process Queue Size (per LPU)
- Quantity of Reductions/Yields/Preempts
- State of the Scheduler (waiting/running)
- Channel State (Blocked/Unblocked)
- ...

◀ ▶ ⏪ ⏩ ⏴ ⏵ ⏶ ⏷ ⏸ ⏹ ⏺ ⏻ ⏼ ⏽ ⏾ ⏿ 🔍 ↺

2014-08-12 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

ErLam Toolkit
Simulation & Visualization
ErLam Toolkit: Simulation & Visualization

ErLam Toolkit: Simulation & Visualization

Things we could log:
■ Process Queue Size (per LPU)
■ Quantity of Reductions/Yields/Preempts
■ State of the Scheduler (waiting/running)
■ Channel State (Blocked/Unblocked)
■ ...

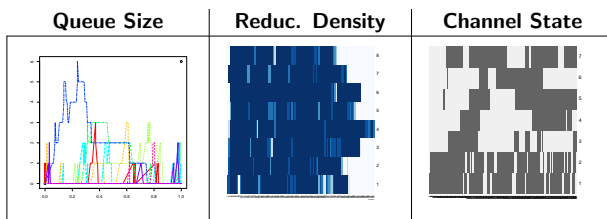
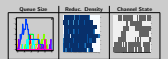
- Queue-Length: work-stealing mechanics and saturation ability.
- Tick-Action: Visualize the density of computation/communication.
- Sched-State: Useful for comparing stealing/process selection mechanics.
- Chan-State: Tracking interactivity, speed of unblock=attentive to cooperation.
- Of course there are more, but we limited ourselves to the above for initial testing purposes.

ErLam Toolkit: Simulation & Visualization

2014-08-12 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

ErLam Toolkit
Simulation & Visualization
ErLam Toolkit: Simulation & Visualization

ErLam Toolkit: Simulation & Visualization



◀ ▶ ⏪ ⏩ ⏴ ⏵ ⏶ ⏷ ⏸ ⏹ ⏺ ⏻ ⏼ ⏽ ⏾ ⏿ 🔍 ↺

- Three types of graphs:
 - Queue Size: X-axis is time, Y-axis is size of queue
 - Density charts: Darkness of the line represents fraction of ticks event happened in.
 - Channel State: dark=blocked, light=unblocked.

Scheduler Implementations

2014-08-12 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Scheduler Implementations
Scheduler Implementations

Scheduler Implementations

■ Scheduler Implementations
■ Example Schedulers
■ Feedback Mechanisms

Introduce the new section:

- 3 Scheduler Implementations
 - Example Schedulers
 - Feedback Mechanisms

3 Scheduler Implementations

- Example Schedulers
- Feedback Mechanisms

◀ ▶ ⏪ ⏩ ⏴ ⏵ ⏶ ⏷ ⏸ ⏹ ⏺ ⏻ ⏼ ⏽ ⏾ ⏿ 🔍 ↺

Scheduler Implementations: Example Schedulers

- (MTRRGQ) - Round-Robin with Single Global Queue
 - All LPU's share a Process Queue.
- (MTRRWS-SQ) - Round-Robin with Work-Stealing via Direct Access
 - All LPU's have their own Process Queue.
 - LPU's can steal processes by grabbing them off the end of another LPU's queue.

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻

Scheduler Implementations: Feedback Mechanisms

Three types of mechanics:

- Longevity-Based Batching
- Channel Pinning
- Bipartite-Graph Aided Sorting

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻

Scheduler Implementations: Feedback Mechanisms

Longevity-Based Batching

- Choose via Round-Robin
 - from batch rather than queue
 - keeps track of number of rounds (batch size)
- Work-Steal whole batches
- Spawn to batch unless: $|b_i| \geq B$
 - Make singleton with new process.
 - Push parent and child into new batch.

◀ ◻ ▶ ◀ ▢ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Scheduler Implementations
 - Example Schedulers
 - Scheduler Implementations: Example Schedulers

- (MTRRSQ) - Round-Robin with Single Global Queue
 - All LPU's share a Process Queue.
- (MTRRSW-SQ) - Round-Robin with Work-Stealing via Direct Access
 - All LPU's have their own Process Queue.
 - LPU's can steal processes by grabbing them off the end of another LPU's queue.

Two of the basic schedulers built where:

1. RR w/ Global Queue: all synchronization around a single shared queue
2. RR w/ Work-Stealing: each scheduler gets their own queue but, they now need to steal work from others.
 - Implemented multiple types of work stealing, but we'll limit talk to one type:
 - Stealing directly from another LPUs by accessing the end of their process queue.

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Scheduler Implementations
 - Feedback Mechanisms
 - Scheduler Implementations: Feedback Mechanisms

Scheduler Implementations: Feedback Mechanisms

These types of mechanism:

- Longevity-Based Batching
- Channel Pinning
- Bipartite-Graph Aided Sorting

- Instead of a single cooperativity-conscious scheduler, we implemented three mechanics which take cooperativity into account on top of the basic schedulers.

- Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages
 - Scheduler Implementations
 - Feedback Mechanisms
 - Scheduler Implementations: Feedback Mechanisms

Scheduler Implementations: Feedback Mechanisms

Complexity Based Scheduling

- Choose via Round-Robin
 - Spin latch rather than queue
 - keeps track of number of rounds (latch size)
- Work-Steal whole latches.
- Spin to latch unless: $|B| \geq B$
 - Make dispatch with new process.
 - Push parent and child into new latch.

- Batching processes based on longevity.
 - Based on $\text{occam-}\Pi$.
 - if a process communicates frequently then it will be batched (absorption), singleton if very computation-bound.
- We are normal RR but with one extra layer.
- If batch is too big during spawns we can:
 - Make singleton, best if child is needed to start work right away. Map-Reduce.
 - Make push-back, parent can get another chance to spawn more children sooner.
- GOAL here is to analyze the effects of grouping the frequently communicating.

Scheduler Implementations: Feedback Mechanisms

- Choose via Round-Robin
 - from batch rather than queue
 - keeps track of number of rounds (batch size)
- Work-Steal whole batches
- Spawn to batch unless: $|b_i| \geq B$
 - Make singleton with new process.
 - Push parent and child into new batch.

GOAL: Can batching based on longevity account for fine/coarse parallelism in application?

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ≡ ≡ ≡ ↺ 🔍 ↻

Scheduler Implementations: Feedback Mechanisms

- Upon call to *newchan*, pin to LPU based on spread algorithm:
 - *same* - LPU *newchan* is called is where it is pinned.
 - *even* - Cycle through LPUs and pin based on that.
 - ...
- Work-steal based on channel that's been pinned to you.

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻

Scheduler Implementations: Feedback Mechanisms

- Upon call to *newchan*, pin to LPU based on spread algorithm:
 - *same* - LPU *newchan* is called is where it is pinned.
 - *even* - Cycle through LPUs and pin based on that.
 - ...
- Work-steal based on channel that's been pinned to you.

GOAL: Can an *even*-like spread increase early saturation?

◀ ◻ ▶ ◀ ▢ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻

12 Process Cooperativity as a Feedback Metric in Concurrent
Message-Passing Languages

- └ Scheduler Implementations

Feedback Mechanisms

- Scheduler Implementations: Feedback Mechanisms

Scheduler Implementations: Feedback Mechanisms
Congruent Round Rerouting

- Choose via Round-Robin
 - from batch rather than queue
 - keeps track of number of rounds (batch size)
- Work-Steal whole batches
- Spawns to batch unless: $|B| \geq B$
 - Make singleton with new process.
 - Push parent and child into new batch.

GOAL: Can batching based on longevity account for file/zone parallelism in application?

- Batching processes based on longevity.
 - Based on $\text{occam-}\Pi$.
 - if a process communicates frequently then it will be batched (absorption), singleton if very computation-bound.
- We are normal RR but with one extra layer.
- If batch is too big during spawns we can:
 - Make singleton, best if child is needed to start work right away. Map-Reduce.
 - Make push-back, parent can get another chance to spawn more children sooner.
- GOAL here is to analyze the effects of grouping the frequently communicating.

12 Message-Passing Languages

Scheduler Implementations

Feedback Mechanisms

- └ Scheduler Implementations: Feedback Mechanisms

- Upon call to *newchan*, pin to LPU based on spread algo:
 - none - LPU *newchan* is called is where it is pinned.
 - even - Cycle through LPUs and pin based on that.
 - ...
- Work-ctrl based on channel that's been pinned to pin.

- Pin channels to LPUs.
 - Pinning a channel means to set a process affinity to a LPU based on the channels it uses.
 - Work-Stealing works like Go-Fish.
- Absorption channels would work well here to relocate processes based on channel usage.
- GOAL: can we improve on work-stealing by being selective.

Bipartite-Graph Aided Sorting

-

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻

2014-08-31

- Feedback Mechanisms

Scheduler Implementations: Feedback Mechanisms

- ### Scheduler Implementations: Feedback Mechanisms
- Round-robin, Round-robin
- Based on Round-Robin & Work-stealing
 - Keep track of events which may affect cooperativity:
 - Spawning
 - Blocking/Unblocking
 - Crash
 - If number of events over some threshold, re-sort
-
- The diagram shows two vertical ovals labeled 'B' and 'C'. Inside oval 'B' are three small circles labeled 'b1', 'b2', and 'b3'. Inside oval 'C' are three small circles labeled 'c1', 'c2', and 'c3'. Arrows indicate interactions: from 'b1' to 'c1', from 'b2' to 'c2', from 'b3' to 'c3', and from 'c1' to 'c2'.

Bipartite-Graph Aided Sorting

-

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻

2014-08-12

-Feedback Mechanisms

-Feedback Mechanisms

Scheduler Implementations: Feedback Mechanisms

- ### Scheduler Implementations: Feedback Mechanisms
- Round-Robin, Round Stealing
- Based on Round-Robin & Round-stealing
 - Keep track of events which may affect cooperativity:
 - Spawning
 - Unlocking/unblocking
 - Block
 - If number of events over some threshold, re-sort
-
- The diagram shows two vertical ovals representing sets of processes, labeled P_1 and P_2 . Inside each oval, there are several small circles representing individual processes. Arrows indicate interactions between processes in P_1 and P_2 . Specifically, there are arrows from the top process in P_1 to the top process in P_2 , and from the bottom process in P_1 to the bottom process in P_2 . There are also arrows from the top process in P_2 to the bottom process in P_1 , and from the bottom process in P_2 to the top process in P_1 , suggesting a bidirectional flow of information or data.
- CSAPP, 4th ed., Chapter 10, Section 10.4, Slide 10-10

GOAL: Are alternate channel implementations worth exploration?

Results

2014-08-12

Results

Results

Introduce the new section:

4 Results

4 Results

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻

Results:

Longevity-Based Batching

- Can batching based on longevity recognize fine/coarse parallelism in an application?

Channel Pinning

- Can an even-like spread increase early saturation?

Bipartite-Graph Aided Sorting

- Are alternate channel implementations worth exploration?

Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12

Results

Results:

- Remind about the goals of the talk.
- LBB: Would like to take advantage of the frequency of communication.

Results

Longevity-Based Batching

- Can batching based on longevity recognize fine/coarse parallelism in an application?

Channel Pinning

- Can an even-like spread increase early saturation?

Bipartite-Graph Aided Sorting

- Are alternate channel implementations worth exploration?

Results: Longevity-Based Batching




Table: Comparison of different sized $PRing_N$ on the Longevity Batching Scheduler with batch size $B = 10$.

Navigation icons: back, forward, search, etc.


Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12


Results

Results: Longevity-Based Batching

- Early tests gave promising results.
- Here is $PRing_N$ which shows the reabsorption and containment on a single LPU as expected and hoped.
- So does batching based on longevity really recognize fine/coarse parallelism in an application?
- Sort of, if you know what the right time-quantum is to make that distinction.



Results:



Navigation icons: back, forward, search, etc.


Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12

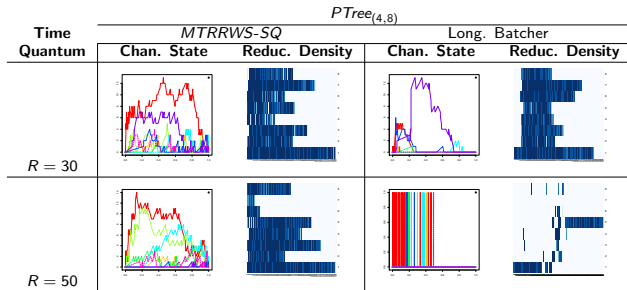
Results

Results:

Reference Slide.



Results: Longevity-Based Batching




Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12


Results

Results: Longevity-Based Batching



- Comparison of $P\text{Tree}_{(4,8)}$ running with the Longevity-Based Batching Scheduler and $M\text{TRRWS-SQ}$ at different time-quantums.
- Absorption channels help here to relocate processes.
- At lower time quantums Long. Batcher starts to look like $M\text{TRRWS-SQ}$, however batching and absorption channels tend to lead to consolidation.
- At higher time quantums Long. Batcher results in the originally expected work-groups. But it turns out to be inefficient due to lost chances of parallelism of each "star" of each group.
- Heuristical adjustment of the time-quantum would definitely be possible.
- NOTE: We don't capture overhead of stealing. Batching has obvious gains here.

Results:




Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12

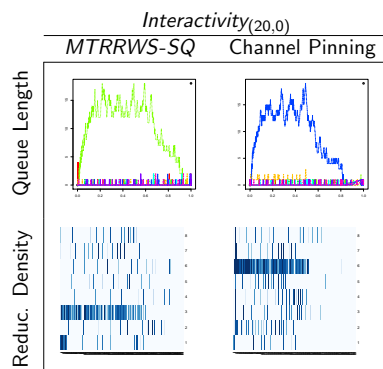
Results

Results:



Reference Slide.

Results: Channel Pinning




Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12


Results

Results: Channel Pinning



- Comparison of Uniform synchronization for $M\text{TRRWS-SQ}$ and the Channel Pinning Scheduler on Absorption Channels.
- This used the *even* spread type.
- Note the speed at which it saturates all cores.
- Despite Naive WS, we still have decent spread.

Results:




Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Results


Results:



Reference Slide.

Results:

Results: Bipartite-Aided Graph Sorting



Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Results

Results:



- There is an issue with our next test.
- Our base primitive behaviours don't have a preferred order of execution.
- Ring does, but there is nothing parallel about it.

Results:

Results: Bipartite-Aided Graph Sorting

```
fun N.
  (omega fun f,m.(
    if (leq m 1)
      m
      (merge fun _.(f f (sub m 1))
              fun _.(f f (sub m 2))
              add)) N)
```

Navigation icons: back, forward, search, etc.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages


Results

Results:



- We therefore used a parallel Fibonacci program.
- Merge function waits for each sub process to finish before running Add.
- If we sort based on who hasn't had a chance to communicate yet, we can preferentialize these MapReduce style applications, while not decrementing possible execution in parents.

Results: Bipartite-Aided Graph Sorting




Navigation icons: back, forward, search, etc.

2014-08-12 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Results

Results: Bipartite-Aided Graph Sorting



- PFib has a strong reliance on order of execution.
- Channel State comparison of Parallel Fibonacci executed on *MTRRWS-SQ* and the Bipartite-Graph Aided Sorting Scheduler.
- Note the large reduction in number of ticks.

Conclusions & Future Work

2014-08-12 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Conclusions & Future Work

Conclusions & Future Work



Introduce the new section:

- 5 Conclusions & Future Work
 - ErLam Toolkit
 - Cooperative Schedulers
 - Cooperativity as a Metric

5 Conclusions & Future Work

- ErLam Toolkit
- Cooperative Schedulers
- Cooperativity as a Metric

Navigation icons: back, forward, search, etc.


Conclusions & Future Work: ErLam Toolkit

2014-08-12 Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

Conclusions & Future Work

ErLam Toolkit

Conclusions & Future Work: ErLam Toolkit



- Test Primitives were nicely composable process behaviours.
 - More research into generating behaviours.
 - More compositions: PTree with Rings.
- Log generation, lots of overhead, but good observations.

- Overall pleased with simulator and achieved its goal.
- Future Work:
 - Generate more test primitives, a good library of them would be nice.
 - Compose them easier and more frequently. Perhaps generating work groups as PRing might have made a better comparison than Interactivity.
 - Clean up log generation (reduce overhead).
 - Process evaluation uses alpha-reduction (can be sped up substantially).
 - Make schedulers more adjustable (different spawn/yields/etc).
 - More Channel implementations

Navigation icons: back, forward, search, etc.

Conclusions & Future Work: Cooperative Schedulers

- **Longevity Batching:**
 - Would benefit from heuristic based Quantum selection.
 - As it stands, limited gain from longevity recognition.
- **Channel Pinning:**
 - Promising saturation and work-stealing mechanic.
- **Bipartite-Graph Aided Sorting:**
 - Supprising results on MapReduce style applications.
 - Worth studying Blocking-Channels further for gains from sorting.

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ≡ ≡ ↺ 🔍 ↻

Conclusions & Future Work: Cooperativity as a Metric

- Possible to recognize and benefit from.
- The three example mechanics are promising and can be extended for practice modern languages, despite simplistic simulation language.
- More to explore:
 - Alternate Message-Passing Types (Asymmetric?)
 - ...

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12

- Conclusions & Future Work
 - Cooperative Schedulers

- Cooperative Schedulers

Conclusions & Future Work

Conclusions & Future Work: Cooperative Schedulers

Conclusions & Future Work: Cooperative Scheduling

- **Longevity Blatching:**
 - Would benefit from heuristic based Quantum solutions.
 - As it stands, limited gain from longevity magnification.
- **Channel Pinning:**
 - Promising saturation and work-stealing mechanic.
- **Bipartite-Graph Aided Sorting:**
 - Suprising results on MapReduce style applications.
 - Worth studying blocking Channels further for gains from

- SS: Would benefit from merging with Channel Pinning. Increase likelihood that sorting puts channel partners close.
- LBB + CP might be interesting as channels could own batches.
- Overhead of Sorting? If implemented in a practice language, would it be worth it? Seems counter-intuitive but promising as is.

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12

Conclusions & Future Work

- Cooperativity as a Metric

Conclusions & Future Work: Cooperativity as a Metric

Conclusions: *E. Coli* Water Contamination in a Man...

- Possible to recognize and benefit from.
- The three example mechanics are promising and can be extended for practical modern languages, despite simple simulation language.
- More to explore:
 - Alternate Message-Passing Types (Asymmetric?)
 - ...

- Promising scheduling mechanics
- More research is necessary in the message-passing implementation for it to be practical in common languages.
- (Asymmetric/Directionality would be first on the list)

Process Cooperativity as a Feedback Metric in Concurrent Message-Passing Languages

2014-08-12

Questions

Questions/Comments?

Questions/Comments?

◀ ◻ ▶ ◀ ▢ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻

Thank You!

- ◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻